

(continued from previous page)

```

.....:          D=lambdax: x['A'] + x['C'])
.....:
Out [80]:
   A  B  C  D
0  1  4  5  6
1  2  5  7  9
2  3  6  9 12

```

In the second expression, `x['C']` will refer to the newly created column, that's equal to `dfa['A'] + dfa['B']`.

Indexing / selection

The basics of indexing are as follows:

Operation	Syntax	Result
Select column	<code>df[col]</code>	Series
Select row by label	<code>df.loc[label]</code>	Series
Select row by integer location	<code>df.iloc[loc]</code>	Series
Slice rows	<code>df[5:10]</code>	DataFrame
Select rows by boolean vector	<code>df[bool_vec]</code>	DataFrame

Row selection, for example, returns a Series whose index is the columns of the DataFrame:

```

In [81]: df.loc['b']
Out [81]:
one          2
bar          2
flag        False
foo          bar
one_trunc    2
Name: b, dtype: object

In [82]: df.iloc[2]
Out [82]:
one          3
bar          3
flag         True
foo          bar
one_trunc    NaN
Name: c, dtype: object

```

For a more exhaustive treatment of sophisticated label-based indexing and slicing, see the [section on indexing](#). We will address the fundamentals of reindexing / conforming to new sets of labels in the [section on reindexing](#).

Data alignment and arithmetic

Data alignment between DataFrame objects automatically align on **both the columns and the index (row labels)**. Again, the resulting object will have the union of the column and row labels.

```
In [83]: df = pd.DataFrame(np.random.randn(10, 4), columns=['A', 'B', 'C', 'D'])
```

```
In [84]: df2 = pd.DataFrame(np.random.randn(7, 3), columns=['A', 'B', 'C'])
```

```
In [85]: df + df2
```

```
Out[85]:
```

	A	B	C	D
0	0.045691	-0.014138	1.380871	NaN
1	-0.955398	-1.501007	0.037181	NaN
2	-0.662690	1.534833	-0.859691	NaN
3	-2.452949	1.237274	-0.133712	NaN
4	1.414490	1.951676	-2.320422	NaN
5	-0.494922	-1.649727	-1.084601	NaN
6	-1.047551	-0.748572	-0.805479	NaN
7	NaN	NaN	NaN	NaN
8	NaN	NaN	NaN	NaN
9	NaN	NaN	NaN	NaN

When doing an operation between DataFrame and Series, the default behavior is to align the Series **index** on the DataFrame **columns**, thus **broadcasting** row-wise. For example:

```
In [86]: df - df.iloc[0]
```

```
Out[86]:
```

	A	B	C	D
0	0.000000	0.000000	0.000000	0.000000
1	-1.359261	-0.248717	-0.453372	-1.754659
2	0.253128	0.829678	0.010026	-1.991234
3	-1.311128	0.054325	-1.724913	-1.620544
4	0.573025	1.500742	-0.676070	1.367331
5	-1.741248	0.781993	-1.241620	-2.053136
6	-1.240774	-0.869551	-0.153282	0.000430
7	-0.743894	0.411013	-0.929563	-0.282386
8	-1.194921	1.320690	0.238224	-1.482644
9	2.293786	1.856228	0.773289	-1.446531

In the special case of working with time series data, if the DataFrame index contains dates, the broadcasting will be column-wise:

```
In [87]: index = pd.date_range('1/1/2000', periods=8)
```

```
In [88]: df = pd.DataFrame(np.random.randn(8, 3), index=index, columns=list('ABC'))
```

```
In [89]: df
```

```
Out[89]:
```

	A	B	C
2000-01-01	-1.226825	0.769804	-1.281247
2000-01-02	-0.727707	-0.121306	-0.097883
2000-01-03	0.695775	0.341734	0.959726
2000-01-04	-1.110336	-0.619976	0.149748
2000-01-05	-0.732339	0.687738	0.176444
2000-01-06	0.403310	-0.154951	0.301624
2000-01-07	-2.179861	-1.369849	-0.954208

(continues on next page)

(continued from previous page)

```

2000-01-08    1.462696 -1.743161 -0.826591

In [90]: type(df['A'])
Out[90]: pandas.core.series.Series

In [91]: df - df['A']
Out[91]:
           2000-01-01 00:00:00  2000-01-02 00:00:00  2000-01-03 00:00:00  2000-01-04
↪00:00:00  ...  2000-01-08 00:00:00    A      B      C
2000-01-01    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-02    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-03    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-04    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-05    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-06    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-07    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN
2000-01-08    ...              NaN  NaN  NaN  NaN  NaN      NaN
↪      NaN  ...              NaN  NaN  NaN  NaN  NaN      NaN

[8 rows x 11 columns]

```

Warning:

```
df - df['A']
```

is now deprecated and will be removed in a future release. The preferred way to replicate this behavior is

```
df.sub(df['A'], axis=0)
```

For explicit control over the matching and broadcasting behavior, see the section on *flexible binary operations*.

Operations with scalars are just as you would expect:

```

In [92]: df * 5 + 2
Out[92]:
           A           B           C
2000-01-01 -4.134126  5.849018 -4.406237
2000-01-02 -1.638535  1.393469  1.510587
2000-01-03  5.478873  3.708672  6.798628
2000-01-04 -3.551681 -1.099880  2.748742
2000-01-05 -1.661697  5.438692  2.882222
2000-01-06  4.016548  1.225246  3.508122
2000-01-07 -8.899303 -4.849247 -2.771039
2000-01-08  9.313480 -6.715805 -2.132955

In [93]: 1 / df
Out[93]:
           A           B           C

```

(continues on next page)

(continued from previous page)

```

2000-01-01 -0.815112  1.299033  -0.780489
2000-01-02 -1.374179 -8.243600 -10.216313
2000-01-03  1.437247  2.926250  1.041965
2000-01-04 -0.900628 -1.612966  6.677871
2000-01-05 -1.365487  1.454041  5.667510
2000-01-06  2.479485 -6.453662  3.315381
2000-01-07 -0.458745 -0.730007 -1.047990
2000-01-08  0.683669 -0.573671 -1.209788

```

```
In [94]: df ** 4
```

```
Out[94]:
```

```

           A           B           C
2000-01-01  2.265327  0.351172  2.694833
2000-01-02  0.280431  0.000217  0.000092
2000-01-03  0.234355  0.013638  0.848376
2000-01-04  1.519910  0.147740  0.000503
2000-01-05  0.287640  0.223714  0.000969
2000-01-06  0.026458  0.000576  0.008277
2000-01-07 22.579530  3.521204  0.829033
2000-01-08  4.577374  9.233151  0.466834

```

Boolean operators work as well:

```
In [95]: df1 = pd.DataFrame({'a': [1, 0, 1], 'b': [0, 1, 1]}, dtype=bool)
```

```
In [96]: df2 = pd.DataFrame({'a': [0, 1, 1], 'b': [1, 1, 0]}, dtype=bool)
```

```
In [97]: df1 & df2
```

```
Out[97]:
```

```

   a      b
0  False False
1  False  True
2   True False

```

```
In [98]: df1 | df2
```

```
Out[98]:
```

```

   a      b
0  True  True
1  True  True
2  True  True

```

```
In [99]: df1 ^ df2
```

```
Out[99]:
```

```

   a      b
0  True  True
1  True False
2 False  True

```

```
In [100]: ~df1
```

```
Out[100]:
```

```

   a      b
0 False  True
1  True False
2 False False

```

Transposing

To transpose, access the `T` attribute (also the `transpose` function), similar to an `ndarray`:

```
# only show the first 5 rows
In [101]: df[:5].T
Out[101]:
      2000-01-01  2000-01-02  2000-01-03  2000-01-04  2000-01-05
A    -1.226825   -0.727707    0.695775   -1.110336   -0.732339
B     0.769804   -0.121306    0.341734   -0.619976    0.687738
C    -1.281247   -0.097883    0.959726    0.149748    0.176444
```

DataFrame interoperability with NumPy functions

Elementwise NumPy ufuncs (`log`, `exp`, `sqrt`, ...) and various other NumPy functions can be used with no issues on `Series` and `DataFrame`, assuming the data within are numeric:

```
In [102]: np.exp(df)
Out[102]:
      A      B      C
2000-01-01  0.293222  2.159342  0.277691
2000-01-02  0.483015  0.885763  0.906755
2000-01-03  2.005262  1.407386  2.610980
2000-01-04  0.329448  0.537957  1.161542
2000-01-05  0.480783  1.989212  1.192968
2000-01-06  1.496770  0.856457  1.352053
2000-01-07  0.113057  0.254145  0.385117
2000-01-08  4.317584  0.174966  0.437538

In [103]: np.asarray(df)
Out[103]:
array([[ -1.2268,   0.7698,  -1.2812],
       [ -0.7277,  -0.1213,  -0.0979],
       [  0.6958,   0.3417,   0.9597],
       [ -1.1103,  -0.62   ,   0.1497],
       [ -0.7323,   0.6877,   0.1764],
       [  0.4033,  -0.155  ,   0.3016],
       [ -2.1799,  -1.3698,  -0.9542],
       [  1.4627,  -1.7432,  -0.8266]])
```

`DataFrame` is not intended to be a drop-in replacement for `ndarray` as its indexing semantics and data model are quite different in places from an `n`-dimensional array.

`Series` implements `__array_ufunc__`, which allows it to work with NumPy's [universal functions](#).

The ufunc is applied to the underlying array in a `Series`.

```
In [104]: ser = pd.Series([1, 2, 3, 4])

In [105]: np.exp(ser)
Out[105]:
0      2.718282
1      7.389056
2     20.085537
3     54.598150
dtype: float64
```

Changed in version 0.25.0: When multiple `Series` are passed to a ufunc, they are aligned before performing the operation.

Like other parts of the library, pandas will automatically align labeled inputs as part of a ufunc with multiple inputs. For example, using `numpy.remainder()` on two `Series` with differently ordered labels will align before the operation.

```
In [106]: ser1 = pd.Series([1, 2, 3], index=['a', 'b', 'c'])
In [107]: ser2 = pd.Series([1, 3, 5], index=['b', 'a', 'c'])

In [108]: ser1
Out[108]:
a    1
b    2
c    3
dtype: int64

In [109]: ser2
Out[109]:
b    1
a    3
c    5
dtype: int64

In [110]: np.remainder(ser1, ser2)
Out[110]:
a    1
b    0
c    3
dtype: int64
```

As usual, the union of the two indices is taken, and non-overlapping values are filled with missing values.

```
In [111]: ser3 = pd.Series([2, 4, 6], index=['b', 'c', 'd'])

In [112]: ser3
Out[112]:
b    2
c    4
d    6
dtype: int64

In [113]: np.remainder(ser1, ser3)
Out[113]:
a    NaN
b    0.0
c    3.0
d    NaN
dtype: float64
```

When a binary ufunc is applied to a `Series` and `Index`, the `Series` implementation takes precedence and a `Series` is returned.

```
In [114]: ser = pd.Series([1, 2, 3])
In [115]: idx = pd.Index([4, 5, 6])
```

(continues on next page)

(continued from previous page)

```
In [116]: np.maximum(ser, idx)
Out[116]:
0      4
1      5
2      6
dtype: int64
```

NumPy ufuncs are safe to apply to *Series* backed by non-ndarray arrays, for example `arrays.SparseArray` (see *Sparse calculation*). If possible, the ufunc is applied without converting the underlying data to an ndarray.

Console display

Very large DataFrames will be truncated to display them in the console. You can also get a summary using `info()`. (Here I am reading a CSV version of the **baseball** dataset from the **plyr** R package):

```
In [117]: baseball = pd.read_csv('data/baseball.csv')

In [118]: print(baseball)
      id  player  year  stint team lg   g  ab  r   h  X2b  X3b  hr   rbi  sb_
→ cs  bb      so  ibb  hbp   sh  sf  gidp
0  88641  womacto01  2006     2  CHN  NL  19  50  6  14   1   0   1   2.0  1.0_
→ 1.0   4   4.0  0.0  0.0  3.0  0.0  0.0
1  88643  schilcu01  2006     1  BOS  AL  31   2  0   1   0   0   0   0.0  0.0_
→ 0.0   0   1.0  0.0  0.0  0.0  0.0  0.0
..   ...      ...   ...   ...   ...   ...  ..   ...  ..   ...   ...   ..   ...   ..._
→ ...   ..   ...   ...   ...   ...   ...   ...
98  89533  aloumo01  2007     1  NYN  NL  87  328  51  112   19   1  13  49.0  3.0_
→ 0.0  27  30.0  5.0  2.0  0.0  3.0  13.0
99  89534  alomasa02  2007     1  NYN  NL   8   22   1   3   1   0   0   0.0  0.0_
→ 0.0   0   3.0  0.0  0.0  0.0  0.0  0.0

[100 rows x 23 columns]

In [119]: baseball.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 23 columns):
#   Column  Non-Null Count  Dtype
---  ---
0    id      100 non-null    int64
1  player  100 non-null    object
2   year    100 non-null    int64
3  stint    100 non-null    int64
4   team    100 non-null    object
5    lg      100 non-null    object
6    g       100 non-null    int64
7   ab      100 non-null    int64
8    r       100 non-null    int64
9    h       100 non-null    int64
10  X2b      100 non-null    int64
11  X3b      100 non-null    int64
12  hr       100 non-null    int64
13  rbi      100 non-null    float64
14  sb       100 non-null    float64
15  cs       100 non-null    float64
```

(continues on next page)

(continued from previous page)

```

16  bb      100 non-null    int64
17  so      100 non-null    float64
18  ibb     100 non-null    float64
19  hbp     100 non-null    float64
20  sh      100 non-null    float64
21  sf      100 non-null    float64
22  gidp    100 non-null    float64
dtypes: float64(9), int64(11), object(3)
memory usage: 18.1+ KB

```

However, using `to_string` will return a string representation of the DataFrame in tabular form, though it won't always fit the console width:

```

In [120]: print(baseball.iloc[-20:, :12].to_string())

```

	id	player	year	stint	team	lg	g	ab	r	h	X2b	X3b
80	89474	finlest01	2007	1	COL	NL	43	94	9	17	3	0
81	89480	embreal01	2007	1	OAK	AL	4	0	0	0	0	0
82	89481	edmonji01	2007	1	SLN	NL	117	365	39	92	15	2
83	89482	easleda01	2007	1	NYN	NL	76	193	24	54	6	0
84	89489	delgaca01	2007	1	NYN	NL	139	538	71	139	30	0
85	89493	cormirh01	2007	1	CIN	NL	6	0	0	0	0	0
86	89494	coninje01	2007	2	NYN	NL	21	41	2	8	2	0
87	89495	coninje01	2007	1	CIN	NL	80	215	23	57	11	1
88	89497	clemero02	2007	1	NYA	AL	2	2	0	1	0	0
89	89498	claytro01	2007	2	BOS	AL	8	6	1	0	0	0
90	89499	claytro01	2007	1	TOR	AL	69	189	23	48	14	0
91	89501	cirilje01	2007	2	ARI	NL	28	40	6	8	4	0
92	89502	cirilje01	2007	1	MIN	AL	50	153	18	40	9	2
93	89521	bondsba01	2007	1	SFN	NL	126	340	75	94	14	0
94	89523	biggicr01	2007	1	HOU	NL	141	517	68	130	31	3
95	89525	benitar01	2007	2	FLO	NL	34	0	0	0	0	0
96	89526	benitar01	2007	1	SFN	NL	19	0	0	0	0	0
97	89530	ausmubr01	2007	1	HOU	NL	117	349	38	82	16	3
98	89533	aloumo01	2007	1	NYN	NL	87	328	51	112	19	1
99	89534	alomasa02	2007	1	NYN	NL	8	22	1	3	1	0

Wide DataFrames will be printed across multiple rows by default:

```

In [121]: pd.DataFrame(np.random.randn(3, 12))
Out[121]:

```

	0	1	2	3	4	5	6	7	8	9	10	11
0	-0.345352	1.314232	0.690579	0.995761	2.396780	0.014871	3.357427	-0.317441	-1.236269	0.896171	-0.487602	-0.082240
1	-2.182937	0.380396	0.084844	0.432390	1.519970	-0.493662	0.600178	0.274230	0.132885	-0.023688	2.410179	1.450520
2	0.206053	-0.251905	-2.213588	1.063327	1.266143	0.299368	-0.863838	0.408204	-1.048089	-0.025747	-0.988387	0.094055

You can change how much to print on a single row by setting the `display.width` option:

```

In [122]: pd.set_option('display.width', 40) # default is 80

In [123]: pd.DataFrame(np.random.randn(3, 12))
Out[123]:

```

	0	1	2	3	4	5	6	7	8	9	10	11
0	-0.345352	1.314232	0.690579	0.995761	2.396780	0.014871	3.357427	-0.317441	-1.236269	0.896171	-0.487602	-0.082240
1	-2.182937	0.380396	0.084844	0.432390	1.519970	-0.493662	0.600178	0.274230	0.132885	-0.023688	2.410179	1.450520
2	0.206053	-0.251905	-2.213588	1.063327	1.266143	0.299368	-0.863838	0.408204	-1.048089	-0.025747	-0.988387	0.094055

(continues on next page)

(continued from previous page)

```
0  1.262731  1.289997  0.082423 -0.055758  0.536580 -0.489682  0.369374 -0.034571 -2.
↪484478 -0.281461  0.030711  0.109121
1  1.126203 -0.977349  1.474071 -0.064034 -1.282782  0.781836 -1.071357  0.441153  2.
↪353925  0.583787  0.221471 -0.744471
2  0.758527  1.729689 -0.964980 -0.845696 -1.340896  1.846883 -1.328865  1.682706 -1.
↪717693  0.888782  0.228440  0.901805
```

You can adjust the max width of the individual columns by setting `display.max_colwidth`

```
In [124]: datafile = {'filename': ['filename_01', 'filename_02'],
.....:                'path': ["media/user_name/storage/folder_01/filename_01",
.....:                        "media/user_name/storage/folder_02/filename_02"]}
.....:

In [125]: pd.set_option('display.max_colwidth', 30)

In [126]: pd.DataFrame(datafile)
Out[126]:
   filename                                path
0  filename_01  media/user_name/storage/fo...
1  filename_02  media/user_name/storage/fo...

In [127]: pd.set_option('display.max_colwidth', 100)

In [128]: pd.DataFrame(datafile)
Out[128]:
   filename                                path
0  filename_01  media/user_name/storage/folder_01/filename_01
1  filename_02  media/user_name/storage/folder_02/filename_02
```

You can also disable this feature via the `expand_frame_repr` option. This will print the table in one block.

DataFrame column attribute access and IPython completion

If a DataFrame column label is a valid Python variable name, the column can be accessed like an attribute:

```
In [129]: df = pd.DataFrame({'foo1': np.random.randn(5),
.....:                      'foo2': np.random.randn(5)})
.....:

In [130]: df
Out[130]:
   foo1    foo2
0  1.171216 -0.858447
1  0.520260  0.306996
2 -1.197071 -0.028665
3 -1.066969  0.384316
4 -0.303421  1.574159

In [131]: df.foo1
Out[131]:
0    1.171216
1     0.520260
2    -1.197071
3    -1.066969
```

(continues on next page)

(continued from previous page)

```
4    -0.303421
Name: fool, dtype: float64
```

The columns are also connected to the [IPython](#) completion mechanism so they can be tab-completed:

```
In [5]: df.fo<TAB> # noqa: E225, E999
df.fool  df.foo2
```

1.4.7 Comparison with other tools

Comparison with R / R libraries

Since pandas aims to provide a lot of the data manipulation and analysis functionality that people use R for, this page was started to provide a more detailed look at the [R language](#) and its many third party libraries as they relate to pandas. In comparisons with R and CRAN libraries, we care about the following things:

- **Functionality / flexibility:** what can/cannot be done with each tool
- **Performance:** how fast are operations. Hard numbers/benchmarks are preferable
- **Ease-of-use:** Is one tool easier/harder to use (you may have to be the judge of this, given side-by-side code comparisons)

This page is also here to offer a bit of a translation guide for users of these R packages.

For transfer of DataFrame objects from pandas to R, one option is to use HDF5 files, see [External compatibility](#) for an example.

Quick reference

We'll start off with a quick reference guide pairing some common R operations using [dplyr](#) with pandas equivalents.

Querying, filtering, sampling

R	pandas
<code>dim(df)</code>	<code>df.shape</code>
<code>head(df)</code>	<code>df.head()</code>
<code>slice(df, 1:10)</code>	<code>df.iloc[:9]</code>
<code>filter(df, col1 == 1, col2 == 1)</code>	<code>df.query('col1 == 1 & col2 == 1')</code>
<code>df[df\$col1 == 1 & df\$col2 == 1,]</code>	<code>df[(df.col1 == 1) & (df.col2 == 1)]</code>
<code>select(df, col1, col2)</code>	<code>df[['col1', 'col2']]</code>
<code>select(df, col1:col3)</code>	<code>df.loc[:, 'col1':'col3']</code>
<code>select(df, -(col1:col3))</code>	<code>df.drop(cols_to_drop, axis=1)</code> but see ¹
<code>distinct(select(df, col1))</code>	<code>df[['col1']].drop_duplicates()</code>
<code>distinct(select(df, col1, col2))</code>	<code>df[['col1', 'col2']].drop_duplicates()</code>
<code>sample_n(df, 10)</code>	<code>df.sample(n=10)</code>
<code>sample_frac(df, 0.01)</code>	<code>df.sample(frac=0.01)</code>

¹ R's shorthand for a subrange of columns (`select(df, col1:col3)`) can be approached cleanly in pandas, if you have the list of columns, for example `df[cols[1:3]]` or `df.drop(cols[1:3])`, but doing this by column name is a bit messy.

Sorting

R	pandas
<code>arrange(df, col1, col2)</code>	<code>df.sort_values(['col1', 'col2'])</code>
<code>arrange(df, desc(col1))</code>	<code>df.sort_values('col1', ascending=False)</code>

Transforming

R	pandas
<code>select(df, col_one = col1)</code>	<code>df.rename(columns={'col1': 'col_one'})['col_one']</code>
<code>rename(df, col_one = col1)</code>	<code>df.rename(columns={'col1': 'col_one'})</code>
<code>mutate(df, c=a-b)</code>	<code>df.assign(c=df['a']-df['b'])</code>

Grouping and summarizing

R	pandas
<code>summary(df)</code>	<code>df.describe()</code>
<code>gdf <- group_by(df, col1)</code>	<code>gdf = df.groupby('col1')</code>
<code>summarise(gdf, avg=mean(col1, na.rm=TRUE))</code>	<code>df.groupby('col1').agg({'col1': 'mean'})</code>
<code>summarise(gdf, total=sum(col1))</code>	<code>df.groupby('col1').sum()</code>

Base R

Slicing with R's `c`

R makes it easy to access `data.frame` columns by name

```
df <- data.frame(a=rnorm(5), b=rnorm(5), c=rnorm(5), d=rnorm(5), e=rnorm(5))
df[, c("a", "c", "e")]
```

or by integer location

```
df <- data.frame(matrix(rnorm(1000), ncol=100))
df[, c(1:10, 25:30, 40, 50:100)]
```

Selecting multiple columns by name in pandas is straightforward

```
In [1]: df = pd.DataFrame(np.random.randn(10, 3), columns=list('abc'))

In [2]: df[['a', 'c']]
Out[2]:
```

	a	c
0	0.469112	-1.509059
1	-1.135632	-0.173215
2	0.119209	-0.861849

(continues on next page)

(continued from previous page)

```

3 -2.104569  1.071804
4  0.721555 -1.039575
5  0.271860  0.567020
6  0.276232 -0.673690
7  0.113648  0.524988
8  0.404705 -1.715002
9 -1.039268 -1.157892

```

```
In [3]: df.loc[:, ['a', 'c']]
```

```
Out [3]:
```

```

      a      c
0  0.469112 -1.509059
1 -1.135632 -0.173215
2  0.119209 -0.861849
3 -2.104569  1.071804
4  0.721555 -1.039575
5  0.271860  0.567020
6  0.276232 -0.673690
7  0.113648  0.524988
8  0.404705 -1.715002
9 -1.039268 -1.157892

```

Selecting multiple noncontiguous columns by integer location can be achieved with a combination of the `iloc` indexer attribute and `numpy.r_`.

```
In [4]: named = list('abcdefg')
```

```
In [5]: n = 30
```

```
In [6]: columns = named + np.arange(len(named), n).tolist()
```

```
In [7]: df = pd.DataFrame(np.random.randn(n, n), columns=columns)
```

```
In [8]: df.iloc[:, np.r_[:10, 24:30]]
```

```
Out [8]:
```

```

      a      b      c      d      e      f      g  ...
↪9      24      25      26      27      28      29
0 -1.344312  0.844885  1.075770 -0.109050  1.643563 -1.469388  0.357021  ... -0.
↪968914 -1.170299 -0.226169  0.410835  0.813850  0.132003 -0.827317
1 -0.076467 -1.187678  1.130127 -1.436737 -1.413681  1.607920  1.024180  ... -2.
↪211372  0.959726 -1.110336 -0.619976  0.149748 -0.732339  0.687738
2  0.176444  0.403310 -0.154951  0.301624 -2.179861 -1.369849 -0.954208  ... -0.
↪826591  0.084844  0.432390  1.519970 -0.493662  0.600178  0.274230
3  0.132885 -0.023688  2.410179  1.450520  0.206053 -0.251905 -2.213588  ... 0.
↪299368 -2.484478 -0.281461  0.030711  0.109121  1.126203 -0.977349
4  1.474071 -0.064034 -1.282782  0.781836 -1.071357  0.441153  2.353925  ... -0.
↪744471 -1.197071 -1.066969 -0.303421 -0.858447  0.306996 -0.028665
..      ...      ...      ...      ...      ...      ...      ...
↪.      ...      ...      ...      ...      ...      ...
25  1.492125 -0.068190  0.681456  1.221829 -0.434352  1.204815 -0.195612  ... -0.
↪796211  1.944517  0.042344 -0.307904  0.428572  0.880609  0.487645
26  0.725238  0.624607 -0.141185 -0.143948 -0.328162  2.095086 -0.608888  ... -2.
↪513465 -0.846188  1.190624  0.778507  1.008500  1.424017  0.717110
27  1.262419  1.950057  0.301038 -0.933858  0.814946  0.181439 -0.110015  ... 0.
↪307941 -1.341814  0.334281 -0.162227  1.007824  2.826008  1.458383
28 -1.585746 -0.899734  0.921494 -0.211762 -0.059182  0.058308  0.915377  ... -3.
↪060395  0.403620 -0.026602 -0.240481  0.577223 -1.088417  0.326687

```

(continues on next page)

(continued from previous page)

```

29 -0.986248  0.169729 -1.158091  1.019673  0.646039  0.917399 -0.010435  ...  0.
↪869610 -1.209247 -0.671466  0.332872 -2.013086 -1.602549  0.333109

[30 rows x 16 columns]

```

aggregate

In R you may want to split data into subsets and compute the mean for each. Using a data.frame called `df` and splitting it into groups `by1` and `by2`:

```

df <- data.frame(
  v1 = c(1, 3, 5, 7, 8, 3, 5, NA, 4, 5, 7, 9),
  v2 = c(11, 33, 55, 77, 88, 33, 55, NA, 44, 55, 77, 99),
  by1 = c("red", "blue", 1, 2, NA, "big", 1, 2, "red", 1, NA, 12),
  by2 = c("wet", "dry", 99, 95, NA, "damp", 95, 99, "red", 99, NA, NA))
aggregate(x=df[, c("v1", "v2")], by=list(mydf2$by1, mydf2$by2), FUN = mean)

```

The `groupby()` method is similar to base R `aggregate` function.

```

In [9]: df = pd.DataFrame(
...:     {'v1': [1, 3, 5, 7, 8, 3, 5, np.nan, 4, 5, 7, 9],
...:      'v2': [11, 33, 55, 77, 88, 33, 55, np.nan, 44, 55, 77, 99],
...:      'by1': ["red", "blue", 1, 2, np.nan, "big", 1, 2, "red", 1, np.nan, 12],
...:      'by2': ["wet", "dry", 99, 95, np.nan, "damp", 95, 99, "red", 99, np.nan,
...:              np.nan]})
...:

In [10]: g = df.groupby(['by1', 'by2'])

In [11]: g[['v1', 'v2']].mean()
Out[11]:

```

		v1	v2
by1	by2		
1	95	5.0	55.0
	99	5.0	55.0
2	95	7.0	77.0
	99	NaN	NaN
big	damp	3.0	33.0
blue	dry	3.0	33.0
red	red	4.0	44.0
	wet	1.0	11.0

For more details and examples see [the groupby documentation](#).

match / %in%

A common way to select data in R is using `%in%` which is defined using the function `match`. The operator `%in%` is used to return a logical vector indicating if there is a match or not:

```
s <- 0:4
s %in% c(2,4)
```

The `isin()` method is similar to R `%in%` operator:

```
In [12]: s = pd.Series(np.arange(5), dtype=np.float32)

In [13]: s.isin([2, 4])
Out[13]:
0    False
1    False
2     True
3    False
4     True
dtype: bool
```

The `match` function returns a vector of the positions of matches of its first argument in its second:

```
s <- 0:4
match(s, c(2,4))
```

For more details and examples see [the reshaping documentation](#).

tapply

`tapply` is similar to `aggregate`, but data can be in a ragged array, since the subclass sizes are possibly irregular. Using a data.frame called `baseball`, and retrieving information based on the array `team`:

```
baseball <-
  data.frame(team = gl(5, 5,
    labels = paste("Team", LETTERS[1:5])),
    player = sample(letters, 25),
    batting.average = runif(25, .200, .400))

tapply(baseball$batting.average, baseball$team,
  max)
```

In pandas we may use `pivot_table()` method to handle this:

```
In [14]: import random

In [15]: import string

In [16]: baseball = pd.DataFrame(
....:     {'team': ["team %d" % (x + 1) for x in range(5)] * 5,
....:      'player': random.sample(list(string.ascii_lowercase), 25),
....:      'batting avg': np.random.uniform(.200, .400, 25)})
....:

In [17]: baseball.pivot_table(values='batting avg', columns='team', aggfunc=np.max)
Out[17]:
```

(continues on next page)

(continued from previous page)

team	team 1	team 2	team 3	team 4	team 5
batting avg	0.352134	0.295327	0.397191	0.394457	0.396194

For more details and examples see [the reshaping documentation](#).

subset

The `query()` method is similar to the base R `subset` function. In R you might want to get the rows of a data frame where one column's values are less than another column's values:

```
df <- data.frame(a=rnorm(10), b=rnorm(10))
subset(df, a <= b)
df[df$a <= df$b,] # note the comma
```

In pandas, there are a few ways to perform subsetting. You can use `query()` or pass an expression as if it were an index/slice as well as standard boolean indexing:

```
In [18]: df = pd.DataFrame({'a': np.random.randn(10), 'b': np.random.randn(10)})
```

```
In [19]: df.query('a <= b')
```

```
Out [19]:
```

```
      a      b
1  0.174950  0.552887
2 -0.023167  0.148084
3 -0.495291 -0.300218
4 -0.860736  0.197378
5 -1.134146  1.720780
7 -0.290098  0.083515
8  0.238636  0.946550
```

```
In [20]: df[df['a'] <= df['b']]
```

```
Out [20]:
```

```
      a      b
1  0.174950  0.552887
2 -0.023167  0.148084
3 -0.495291 -0.300218
4 -0.860736  0.197378
5 -1.134146  1.720780
7 -0.290098  0.083515
8  0.238636  0.946550
```

```
In [21]: df.loc[df['a'] <= df['b']]
```

```
Out [21]:
```

```
      a      b
1  0.174950  0.552887
2 -0.023167  0.148084
3 -0.495291 -0.300218
4 -0.860736  0.197378
5 -1.134146  1.720780
7 -0.290098  0.083515
8  0.238636  0.946550
```

For more details and examples see [the query documentation](#).

with

An expression using a data.frame called `df` in R with the columns `a` and `b` would be evaluated using `with` like so:

```
df <- data.frame(a=rnorm(10), b=rnorm(10))
with(df, a + b)
df$a + df$b # same as the previous expression
```

In pandas the equivalent expression, using the `eval()` method, would be:

```
In [22]: df = pd.DataFrame({'a': np.random.randn(10), 'b': np.random.randn(10)})

In [23]: df.eval('a + b')
Out[23]:
0    -0.091430
1    -2.483890
2    -0.252728
3    -0.626444
4    -0.261740
5     2.149503
6    -0.332214
7     0.799331
8    -2.377245
9     2.104677
dtype: float64

In [24]: df['a'] + df['b'] # same as the previous expression
Out[24]:
0    -0.091430
1    -2.483890
2    -0.252728
3    -0.626444
4    -0.261740
5     2.149503
6    -0.332214
7     0.799331
8    -2.377245
9     2.104677
dtype: float64
```

In certain cases `eval()` will be much faster than evaluation in pure Python. For more details and examples see [the eval documentation](#).

plyr

`plyr` is an R library for the split-apply-combine strategy for data analysis. The functions revolve around three data structures in R, `a` for arrays, `l` for lists, and `d` for data.frame. The table below shows how these data structures could be mapped in Python.

R	Python
array	list
lists	dictionary or list of objects
data.frame	dataframe

ddply

An expression using a data.frame called `df` in R where you want to summarize `x` by month:

```
require(plyr)
df <- data.frame(
  x = runif(120, 1, 168),
  y = runif(120, 7, 334),
  z = runif(120, 1.7, 20.7),
  month = rep(c(5,6,7,8),30),
  week = sample(1:4, 120, TRUE)
)

ddply(df, .(month, week), summarize,
  mean = round(mean(x), 2),
  sd = round(sd(x), 2))
```

In pandas the equivalent expression, using the `groupby()` method, would be:

```
In [25]: df = pd.DataFrame({'x': np.random.uniform(1., 168., 120),
.....:                    'y': np.random.uniform(7., 334., 120),
.....:                    'z': np.random.uniform(1.7, 20.7, 120),
.....:                    'month': [5, 6, 7, 8] * 30,
.....:                    'week': np.random.randint(1, 4, 120)})
.....:

In [26]: grouped = df.groupby(['month', 'week'])

In [27]: grouped['x'].agg([np.mean, np.std])
Out[27]:
```

		mean	std
month	week		
5	1	63.653367	40.601965
	2	78.126605	53.342400
	3	92.091886	57.630110
6	1	81.747070	54.339218
	2	70.971205	54.687287
	3	100.968344	54.010081
7	1	61.576332	38.844274
	2	61.733510	48.209013
	3	71.688795	37.595638
8	1	62.741922	34.618153
	2	91.774627	49.790202
	3	73.936856	60.773900

For more details and examples see [the groupby documentation](#).

reshape / reshape2

melt.array

An expression using a 3 dimensional array called `a` in R where you want to melt it into a data.frame:

```
a <- array(c(1:23, NA), c(2,3,4))
data.frame(melt(a))
```

In Python, since `a` is a list, you can simply use list comprehension.

```
In [28]: a = np.array(list(range(1, 24)) + [np.NaN]).reshape(2, 3, 4)

In [29]: pd.DataFrame([tuple(list(x) + [val]) for x, val in np.ndenumerate(a)])
Out[29]:
```

	0	1	2	3
0	0	0	0	1.0
1	0	0	1	2.0
2	0	0	2	3.0
3	0	0	3	4.0
4	0	1	0	5.0
..
19	1	1	3	20.0
20	1	2	0	21.0
21	1	2	1	22.0
22	1	2	2	23.0
23	1	2	3	NaN

[24 rows x 4 columns]

melt.list

An expression using a list called `a` in R where you want to melt it into a data.frame:

```
a <- as.list(c(1:4, NA))
data.frame(melt(a))
```

In Python, this list would be a list of tuples, so `DataFrame()` method would convert it to a dataframe as required.

```
In [30]: a = list(enumerate(list(range(1, 5)) + [np.NaN]))

In [31]: pd.DataFrame(a)
Out[31]:
```

	0	1
0	0	1.0
1	1	2.0
2	2	3.0
3	3	4.0
4	4	NaN

For more details and examples see [the Into to Data Structures documentation](#).

melt.data.frame

An expression using a data.frame called `cheese` in R where you want to reshape the data.frame:

```
cheese <- data.frame(  
  first = c('John', 'Mary'),  
  last = c('Doe', 'Bo'),  
  height = c(5.5, 6.0),  
  weight = c(130, 150)  
)  
melt(cheese, id=c("first", "last"))
```

In Python, the `melt()` method is the R equivalent:

```
In [32]: cheese = pd.DataFrame({'first': ['John', 'Mary'],  
  ....:                        'last': ['Doe', 'Bo'],  
  ....:                        'height': [5.5, 6.0],  
  ....:                        'weight': [130, 150]})  
  ....:  
  
In [33]: pd.melt(cheese, id_vars=['first', 'last'])  
Out[33]:  
   first last variable  value  
0  John  Doe   height    5.5  
1  Mary   Bo   height    6.0  
2  John  Doe   weight   130.0  
3  Mary   Bo   weight   150.0  
  
In [34]: cheese.set_index(['first', 'last']).stack() # alternative way  
Out[34]:  
first last  
John  Doe   height    5.5  
       weight   130.0  
Mary   Bo   height    6.0  
       weight   150.0  
dtype: float64
```

For more details and examples see [the reshaping documentation](#).

cast

In R `acast` is an expression using a data.frame called `df` in R to cast into a higher dimensional array:

```
df <- data.frame(  
  x = runif(12, 1, 168),  
  y = runif(12, 7, 334),  
  z = runif(12, 1.7, 20.7),  
  month = rep(c(5,6,7),4),  
  week = rep(c(1,2), 6)  
)  
  
mdf <- melt(df, id=c("month", "week"))  
acast(mdf, week ~ month ~ variable, mean)
```

In Python the best way is to make use of `pivot_table()`:

```
In [35]: df = pd.DataFrame({'x': np.random.uniform(1., 168., 12),
.....:                    'y': np.random.uniform(7., 334., 12),
.....:                    'z': np.random.uniform(1.7, 20.7, 12),
.....:                    'month': [5, 6, 7] * 4,
.....:                    'week': [1, 2] * 6})
.....:

In [36]: mdf = pd.melt(df, id_vars=['month', 'week'])

In [37]: pd.pivot_table(mdf, values='value', index=['variable', 'week'],
.....:                  columns=['month'], aggfunc=np.mean)
.....:
Out[37]:
```

		5	6	7
variable	week			
x	1	93.888747	98.762034	55.219673
	2	94.391427	38.112932	83.942781
y	1	94.306912	279.454811	227.840449
	2	87.392662	193.028166	173.899260
z	1	11.016009	10.079307	16.170549
	2	8.476111	17.638509	19.003494

Similarly for `dcast` which uses a data.frame called `df` in R to aggregate information based on `Animal` and `FeedType`:

```
df <- data.frame(
  Animal = c('Animal1', 'Animal2', 'Animal3', 'Animal2', 'Animal1',
            'Animal2', 'Animal3'),
  FeedType = c('A', 'B', 'A', 'A', 'B', 'B', 'A'),
  Amount = c(10, 7, 4, 2, 5, 6, 2)
)

dcast(df, Animal ~ FeedType, sum, fill=NaN)
# Alternative method using base R
with(df, tapply(Amount, list(Animal, FeedType), sum))
```

Python can approach this in two different ways. Firstly, similar to above using `pivot_table()`:

```
In [38]: df = pd.DataFrame({
.....:   'Animal': ['Animal1', 'Animal2', 'Animal3', 'Animal2', 'Animal1',
.....:            'Animal2', 'Animal3'],
.....:   'FeedType': ['A', 'B', 'A', 'A', 'B', 'B', 'A'],
.....:   'Amount': [10, 7, 4, 2, 5, 6, 2],
.....: })
.....:

In [39]: df.pivot_table(values='Amount', index='Animal', columns='FeedType',
.....:                  aggfunc='sum')
.....:
Out[39]:
```

	A	B
Animal		
Animal1	10.0	5.0
Animal2	2.0	13.0
Animal3	6.0	NaN

The second approach is to use the `groupby()` method:

```
In [40]: df.groupby(['Animal', 'FeedType'])['Amount'].sum()
Out[40]:
Animal  FeedType
Animal1  A          10
         B           5
Animal2  A           2
         B          13
Animal3  A           6
Name: Amount, dtype: int64
```

For more details and examples see [the reshaping documentation](#) or [the groupby documentation](#).

factor

pandas has a data type for categorical data.

```
cut(c(1,2,3,4,5,6), 3)
factor(c(1,2,3,2,2,3))
```

In pandas this is accomplished with `pd.cut` and `astype("category")`:

```
In [41]: pd.cut(pd.Series([1, 2, 3, 4, 5, 6]), 3)
Out[41]:
0    (0.995, 2.667]
1    (0.995, 2.667]
2    (2.667, 4.333]
3    (2.667, 4.333]
4    (4.333, 6.0]
5    (4.333, 6.0]
dtype: category
Categories (3, interval[float64]): [(0.995, 2.667] < (2.667, 4.333] < (4.333, 6.0]]

In [42]: pd.Series([1, 2, 3, 2, 2, 3]).astype("category")
Out[42]:
0    1
1    2
2    3
3    2
4    2
5    3
dtype: category
Categories (3, int64): [1, 2, 3]
```

For more details and examples see [categorical introduction](#) and the [API documentation](#). There is also a documentation regarding the [differences to R's factor](#).

Comparison with SQL

Since many potential pandas users have some familiarity with [SQL](#), this page is meant to provide some examples of how various SQL operations would be performed using pandas.

If you're new to pandas, you might want to first read through [10 Minutes to pandas](#) to familiarize yourself with the library.

As is customary, we import pandas and NumPy as follows:

```
In [1]: import pandas as pd
In [2]: import numpy as np
```

Most of the examples will utilize the `tips` dataset found within pandas tests. We'll read the data into a DataFrame called `tips` and assume we have a database table of the same name and structure.

```
In [3]: url = ('https://raw.githubusercontent.com/pandas-dev/
...:         '/pandas/master/pandas/tests/io/data/csv/tips.csv')
...:

In [4]: tips = pd.read_csv(url)

In [5]: tips.head()
Out[5]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

SELECT

In SQL, selection is done using a comma-separated list of columns you'd like to select (or a `*` to select all columns):

```
SELECT total_bill, tip, smoker, time
FROM tips
LIMIT 5;
```

With pandas, column selection is done by passing a list of column names to your DataFrame:

```
In [6]: tips[['total_bill', 'tip', 'smoker', 'time']].head(5)
Out[6]:
```

	total_bill	tip	smoker	time
0	16.99	1.01	No	Dinner
1	10.34	1.66	No	Dinner
2	21.01	3.50	No	Dinner
3	23.68	3.31	No	Dinner
4	24.59	3.61	No	Dinner

Calling the DataFrame without the list of column names would display all columns (akin to SQL's `*`).

In SQL, you can add a calculated column:

```
SELECT *, tip/total_bill as tip_rate
FROM tips
LIMIT 5;
```

With pandas, you can use the `DataFrame.assign()` method of a `DataFrame` to append a new column:

```
In [7]: tips.assign(tip_rate=tips['tip'] / tips['total_bill']).head(5)
Out[7]:
```

	total_bill	tip	sex	smoker	day	time	size	tip_rate
0	16.99	1.01	Female	No	Sun	Dinner	2	0.059447
1	10.34	1.66	Male	No	Sun	Dinner	3	0.160542
2	21.01	3.50	Male	No	Sun	Dinner	3	0.166587
3	23.68	3.31	Male	No	Sun	Dinner	2	0.139780
4	24.59	3.61	Female	No	Sun	Dinner	4	0.146808

WHERE

Filtering in SQL is done via a `WHERE` clause.

```
SELECT *
FROM tips
WHERE time = 'Dinner'
LIMIT 5;
```

`DataFrames` can be filtered in multiple ways; the most intuitive of which is using `boolean indexing`.

```
In [8]: tips[tips['time'] == 'Dinner'].head(5)
Out[8]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

The above statement is simply passing a `Series` of `True/False` objects to the `DataFrame`, returning all rows with `True`.

```
In [9]: is_dinner = tips['time'] == 'Dinner'

In [10]: is_dinner.value_counts()
Out[10]:
True      176
False      68
Name: time, dtype: int64

In [11]: tips[is_dinner].head(5)
Out[11]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

Just like SQL's `OR` and `AND`, multiple conditions can be passed to a `DataFrame` using `|` (`OR`) and `&` (`AND`).

```
-- tips of more than $5.00 at Dinner meals
SELECT *
FROM tips
WHERE time = 'Dinner' AND tip > 5.00;
```

```
# tips of more than $5.00 at Dinner meals
In [12]: tips[(tips['time'] == 'Dinner') & (tips['tip'] > 5.00)]
Out[12]:
```

	total_bill	tip	sex	smoker	day	time	size
23	39.42	7.58	Male	No	Sat	Dinner	4
44	30.40	5.60	Male	No	Sun	Dinner	4
47	32.40	6.00	Male	No	Sun	Dinner	4
52	34.81	5.20	Female	No	Sun	Dinner	4
59	48.27	6.73	Male	No	Sat	Dinner	4
116	29.93	5.07	Male	No	Sun	Dinner	4
155	29.85	5.14	Female	No	Sun	Dinner	5
170	50.81	10.00	Male	Yes	Sat	Dinner	3
172	7.25	5.15	Male	Yes	Sun	Dinner	2
181	23.33	5.65	Male	Yes	Sun	Dinner	2
183	23.17	6.50	Male	Yes	Sun	Dinner	4
211	25.89	5.16	Male	Yes	Sat	Dinner	4
212	48.33	9.00	Male	No	Sat	Dinner	4
214	28.17	6.50	Female	Yes	Sat	Dinner	3
239	29.03	5.92	Male	No	Sat	Dinner	3

```
-- tips by parties of at least 5 diners OR bill total was more than $45
SELECT *
FROM tips
WHERE size >= 5 OR total_bill > 45;
```

```
# tips by parties of at least 5 diners OR bill total was more than $45
In [13]: tips[(tips['size'] >= 5) | (tips['total_bill'] > 45)]
Out[13]:
```

	total_bill	tip	sex	smoker	day	time	size
59	48.27	6.73	Male	No	Sat	Dinner	4
125	29.80	4.20	Female	No	Thur	Lunch	6
141	34.30	6.70	Male	No	Thur	Lunch	6
142	41.19	5.00	Male	No	Thur	Lunch	5
143	27.05	5.00	Female	No	Thur	Lunch	6
155	29.85	5.14	Female	No	Sun	Dinner	5
156	48.17	5.00	Male	No	Sun	Dinner	6
170	50.81	10.00	Male	Yes	Sat	Dinner	3
182	45.35	3.50	Male	Yes	Sun	Dinner	3
185	20.69	5.00	Male	No	Sun	Dinner	5
187	30.46	2.00	Male	Yes	Sun	Dinner	5
212	48.33	9.00	Male	No	Sat	Dinner	4
216	28.15	3.00	Male	Yes	Sat	Dinner	5

NULL checking is done using the `notna()` and `isna()` methods.

```
In [14]: frame = pd.DataFrame({'col1': ['A', 'B', np.NaN, 'C', 'D'],
.....:                        'col2': ['F', np.NaN, 'G', 'H', 'I']})
.....:

In [15]: frame
Out[15]:
```

(continues on next page)

(continued from previous page)

```

    col1 col2
0      A    F
1      B  NaN
2    NaN    G
3      C    H
4      D    I

```

Assume we have a table of the same structure as our DataFrame above. We can see only the records where `col2` IS NULL with the following query:

```

SELECT *
FROM frame
WHERE col2 IS NULL;

```

```

In [16]: frame[frame['col2'].isna()]
Out[16]:
    col1 col2
1      B  NaN

```

Getting items where `col1` IS NOT NULL can be done with `notna()`.

```

SELECT *
FROM frame
WHERE col1 IS NOT NULL;

```

```

In [17]: frame[frame['col1'].notna()]
Out[17]:
    col1 col2
0      A    F
1      B  NaN
3      C    H
4      D    I

```

GROUP BY

In pandas, SQL's GROUP BY operations are performed using the similarly named `groupby()` method. `groupby()` typically refers to a process where we'd like to split a dataset into groups, apply some function (typically aggregation), and then combine the groups together.

A common SQL operation would be getting the count of records in each group throughout a dataset. For instance, a query getting us the number of tips left by sex:

```

SELECT sex, count(*)
FROM tips
GROUP BY sex;
/*
Female      87
Male       157
*/

```

The pandas equivalent would be:

```

In [18]: tips.groupby('sex').size()
Out[18]:

```

(continues on next page)