```
one -2.213588
baz one 1.063327
        1.266143
bar one
         0.299368
qux two
        -0.863838
baz two
         0.408204
bar two
dtype: float64
In [104]: s.sort_index()
Out[104]:
        1.266143
bar one
    two 0.408204
baz one 1.063327
    two -0.863838
foo one -2.213588
    two -0.251905
        0.206053
qux one
    two 0.299368
dtype: float64
In [105]: s.sort_index(level=0)
Out[105]:
bar one 1.266143
    two 0.408204
baz one 1.063327
    two -0.863838
foo one -2.213588
    two -0.251905
qux one
        0.206053
        0.299368
    t.wo
dtype: float64
In [106]: s.sort_index(level=1)
Out[106]:
        1.266143
bar one
         1.063327
baz one
foo one -2.213588
qux one 0.206053
bar two 0.408204
baz two -0.863838
foo two -0.251905
qux two 0.299368
dtype: float64
```

You may also pass a level name to sort_index if the MultiIndex levels are named.

```
0.206053
qux one
          0.299368
    t.wo
dtype: float64
In [109]: s.sort_index(level='L2')
Out[109]:
L1 L2
bar one
          1.266143
          1.063327
baz one
foo one -2.213588
qux one 0.206053
bar two 0.408204
baz two -0.863838
foo two -0.251905
gux two 0.299368
dtype: float64
```

On higher dimensional objects, you can sort any of the other axes by level if they have a MultiIndex:

Indexing will work even if the data are not sorted, but will be rather inefficient (and show a PerformanceWarning). It will also return a copy of the data rather than a view:

```
In [111]: dfm = pd.DataFrame({'jim': [0, 0, 1, 1],
                                'joe': ['x', 'x', 'z', 'y'],
   . . . . . :
                                'jolie': np.random.rand(4)})
   . . . . . :
   . . . . . :
In [112]: dfm = dfm.set_index(['jim', 'joe'])
In [113]: dfm
Out [113]:
             jolie
jim joe
         0.490671
  X
         0.120248
    X
         0.537020
   Z
         0.110968
    У
```

Furthermore, if you try to index something that is not fully lexsorted, this can raise:

```
In [5]: dfm.loc[(0, 'y'):(1, 'z')]
UnsortedIndexError: 'Key length (2) was greater than MultiIndex lexsort depth (1)'
```

The *is_lexsorted()* method on a MultiIndex shows if the index is sorted, and the lexsort_depth property returns the sort depth:

```
In [114]: dfm.index.is_lexsorted()
Out[114]: False
In [115]: dfm.index.lexsort_depth
Out[115]: 1
```

```
In [116]: dfm = dfm.sort_index()
In [117]: dfm
Out [117]:
            jolie
jim joe
0 x
         0.490671
         0.120248
   Х
         0.110968
   У
         0.537020
In [118]: dfm.index.is_lexsorted()
Out [118]: True
In [119]: dfm.index.lexsort_depth
Out[119]: 2
```

And now selection works as expected.

2.3.4 Take methods

Similar to NumPy ndarrays, pandas Index, Series, and DataFrame also provides the take() method that retrieves elements along a given axis at the given indices. The given indices must be either a list or an ndarray of integer index positions. take will also accept negative integers as relative positions to the end of the object.

```
In [125]: index.take(positions)
Out[125]: Int64Index([214, 329, 567], dtype='int64')
In [126]: ser = pd.Series(np.random.randn(10))
In [127]: ser.iloc[positions]
Out [127]:
   -0.179666
    1.824375
   0.392149
dtype: float64
In [128]: ser.take(positions)
Out [128]:
   -0.179666
9
    1.824375
3
    0.392149
dtype: float64
```

For DataFrames, the given indices should be a 1d list or ndarray that specifies row or column positions.

It is important to note that the take method on pandas objects are not intended to work on boolean indices and may return unexpected results.

```
1 -0.223540

1 -0.223540

dtype: float64

In [137]: ser.iloc[[0, 1]]

Out[137]:

0 0.233141

1 -0.223540

dtype: float64
```

Finally, as a small note on performance, because the take method handles a narrower range of inputs, it can offer performance that is a good deal faster than fancy indexing.

```
In [142]: ser = pd.Series(arr[:, 0])
In [143]: %timeit ser.iloc[indexer]
    ....: %timeit ser.take(indexer)
    ....:
200 us +- 15 us per loop (mean +- std. dev. of 7 runs, 10000 loops each)
169 us +- 2.1 us per loop (mean +- std. dev. of 7 runs, 10000 loops each)
```

2.3.5 Index types

We have discussed MultiIndex in the previous sections pretty extensively. Documentation about DatetimeIndex and PeriodIndex are shown here, and documentation about TimedeltaIndex is found here.

In the following sub-sections we will highlight some other index types.

CategoricalIndex

Categorical Index is a type of index that is useful for supporting indexing with duplicates. This is a container around a Categorical and allows efficient indexing and storage of an index with a large number of duplicated elements.

```
In [146]: df['B'] = df['B'].astype(CategoricalDtype(list('cab')))
In [147]: df
Out [147]:
  A B
0 0 a
  1
  2
3 3 b
4 4 c
5 5 a
In [148]: df.dtypes
Out [148]:
       int64
    category
dtype: object
In [149]: df['B'].cat.categories
Out[149]: Index(['c', 'a', 'b'], dtype='object')
```

Setting the index will create a CategoricalIndex.

Indexing with __getitem__/.iloc/.loc works similarly to an Index with duplicates. The indexers must be in the category or the operation will raise a KeyError.

```
In [152]: df2.loc['a']
Out[152]:
    A
B
a  0
a  1
a  5
```

The CategoricalIndex is preserved after indexing:

Sorting the index will sort by the order of the categories (recall that we created the index with CategoricalDtype(list('cab')), so the sorted order is cab).

```
In [154]: df2.sort_index()
Out[154]:
    A
B
c    4
a    0
a    1
```

```
a 5
b 2
b 3
```

Groupby operations on the index will preserve the index nature as well.

```
In [155]: df2.groupby(level=0).sum()
Out [155]:
    A
B
c    4
a    6
b    5

In [156]: df2.groupby(level=0).sum().index
Out [156]: CategoricalIndex(['c', 'a', 'b'], categories=['c', 'a', 'b'], ordered=False,
    → name='B', dtype='category')
```

Reindexing operations will return a resulting index based on the type of the passed indexer. Passing a list will return a plain-old Index; indexing with a Categorical will return a CategoricalIndex, indexed according to the categories of the passed Categorical dtype. This allows one to arbitrarily index these even with values **not** in the categories, similarly to how you can reindex **any** pandas index.

```
In [160]: df3.reindex(['a', 'e'])
Out[160]:
     Α
В
  0.0
а
  NaN
In [161]: df3.reindex(['a', 'e']).index
Out[161]: Index(['a', 'e'], dtype='object', name='B')
In [162]: df3.reindex(pd.Categorical(['a', 'e'], categories=list('abe')))
Out[162]:
     Α
В
  0.0
а
  NaN
In [163]: df3.reindex(pd.Categorical(['a', 'e'], categories=list('abe'))).index
Out[163]: CategoricalIndex(['a', 'e'], categories=['a', 'b', 'e'], ordered=False,
→name-'B', dtype-'category')
                                                                           (continues on next page)
```

Warning: Reshaping and Comparison operations on a CategoricalIndex must have the same categories or a TypeError will be raised.

```
In [164]: df4 = pd.DataFrame({'A': np.arange(2),
                               'B': list('ba')})
   . . . . . :
   . . . . . :
In [165]: df4['B'] = df4['B'].astype(CategoricalDtype(list('ab')))
In [166]: df4 = df4.set_index('B')
In [167]: df4.index
Out[167]: CategoricalIndex(['b', 'a'], categories=['a', 'b'], ordered=False, name='B
→', dtype='category')
In [168]: df5 = pd.DataFrame({'A': np.arange(2),
                               'B': list('bc')})
   . . . . . :
In [169]: df5['B'] = df5['B'].astype(CategoricalDtype(list('bc')))
In [170]: df5 = df5.set_index('B')
In [171]: df5.index
Out[171]: CategoricalIndex(['b', 'c'], categories=['b', 'c'], ordered=False, name='B
\hookrightarrow', dtype='category')
In [1]: pd.concat([df4, df5])
TypeError: categories must match existing categories when appending
```

Int64Index and RangeIndex

Int 64 Index is a fundamental basic index in pandas. This is an immutable array implementing an ordered, sliceable set.

RangeIndex is a sub-class of Int64Index that provides the default index for all NDFrame objects. RangeIndex is an optimized version of Int64Index that can represent a monotonic ordered set. These are analogous to Python range types.

Float64Index

By default a Float64Index will be automatically created when passing floating, or mixed-integer-floating values in index creation. This enables a pure label-based slicing paradigm that makes [], ix, loc for scalar indexing and slicing work exactly the same.

```
In [172]: indexf = pd.Index([1.5, 2, 3, 4.5, 5])
In [173]: indexf
Out[173]: Float64Index([1.5, 2.0, 3.0, 4.5, 5.0], dtype='float64')
In [174]: sf = pd.Series(range(5), index=indexf)
```

```
In [175]: sf
Out[175]:
1.5     0
2.0     1
3.0     2
4.5     3
5.0     4
dtype: int64
```

Scalar selection for [], .loc will always be label based. An integer will match an equal float index (e.g. 3 is equivalent to 3.0).

```
In [176]: sf[3]
Out[176]: 2

In [177]: sf[3.0]
Out[177]: 2

In [178]: sf.loc[3]
Out[178]: 2

In [179]: sf.loc[3.0]
```

The only positional indexing is via iloc.

```
In [180]: sf.iloc[3]
Out[180]: 3
```

A scalar index that is not found will raise a KeyError. Slicing is primarily on the values of the index when using [], ix, loc, and **always** positional when using iloc. The exception is when the slice is boolean, in which case it will always be positional.

```
In [181]: sf[2:4]
Out[181]:
2.0
    1
3.0
      2
dtype: int64
In [182]: sf.loc[2:4]
Out [182]:
2.0
     1
3.0
      2
dtype: int64
In [183]: sf.iloc[2:4]
Out[183]:
3.0
4.5
       3
dtype: int64
```

In float indexes, slicing using floats is allowed.

```
In [184]: sf[2.1:4.6]
Out[184]:
```

```
3.0 2

4.5 3

dtype: int64

In [185]: sf.loc[2.1:4.6]

Out[185]:

3.0 2

4.5 3

dtype: int64
```

In non-float indexes, slicing using floats will raise a TypeError.

Here is a typical use-case for using this type of indexing. Imagine that you have a somewhat irregular timedelta-like indexing scheme, but the data is recorded as floats. This could, for example, be millisecond offsets.

```
In [186]: dfir = pd.concat([pd.DataFrame(np.random.randn(5, 2),
                                          index=np.arange(5) \star 250.0,
                                          columns=list('AB')),
                             pd.DataFrame(np.random.randn(6, 2),
                                          index=np.arange(4, 10) * 250.1,
   . . . . . :
                                          columns=list('AB'))])
   . . . . . :
   . . . . . :
In [187]: dfir
Out [187]:
      -0.435772 -1.188928
250.0 -0.808286 -0.284634
500.0 -1.815703 1.347213
750.0 -0.243487 0.514704
1000.0 1.162969 -0.287725
1000.4 -0.179734 0.993962
1250.5 -0.212673 0.909872
1500.6 -0.733333 -0.349893
1750.7 0.456434 -0.306735
2000.8 0.553396 0.166221
2250.9 -0.101684 -0.734907
```

Selection operations then will always work on a value basis, for all selection operators.

```
In [188]: dfir[0:1000.4]
Out[188]:

A B

0.0 -0.435772 -1.188928
250.0 -0.808286 -0.284634
500.0 -1.815703 1.347213
750.0 -0.243487 0.514704
1000.0 1.162969 -0.287725
1000.4 -0.179734 0.993962
```

```
In [189]: dfir.loc[0:1001, 'A']
Out[189]:
0.0
        -0.435772
250.0 -0.808286
500.0
        -1.815703
750.0
        -0.243487
        1.162969
1000.0
1000.4 -0.179734
Name: A, dtype: float64
In [190]: dfir.loc[1000.4]
Out [190]:
  -0.179734
   0.993962
Name: 1000.4, dtype: float64
```

You could retrieve the first 1 second (1000 ms) of data as such:

```
In [191]: dfir[0:1000]
Out[191]:

A B

0.0 -0.435772 -1.188928
250.0 -0.808286 -0.284634
500.0 -1.815703 1.347213
750.0 -0.243487 0.514704
1000.0 1.162969 -0.287725
```

If you need integer based selection, you should use iloc:

```
In [192]: dfir.iloc[0:5]
Out[192]:

A B

0.0 -0.435772 -1.188928
250.0 -0.808286 -0.284634
500.0 -1.815703 1.347213
750.0 -0.243487 0.514704
1000.0 1.162969 -0.287725
```

Intervallndex

IntervalIndex together with its own dtype, IntervalDtype as well as the Interval scalar type, allow first-class support in pandas for interval notation.

The IntervalIndex allows some unique indexing and is also used as a return type for the categories in cut() and qcut().

Indexing with an IntervalIndex

An IntervalIndex can be used in Series and in DataFrame as the index.

Label based indexing via .loc along the edges of an interval works as you would expect, selecting that particular interval.

If you select a label *contained* within an interval, this will also select the interval.

Selecting using an Interval will only return exact matches (starting from pandas 0.25.0).

```
In [199]: df.loc[pd.Interval(1, 2)]
Out[199]:
A    2
Name: (1, 2], dtype: int64
```

Trying to select an Interval that is not exactly contained in the IntervalIndex will raise a KeyError.

Selecting all Intervals that overlap a given Interval can be performed using the *overlaps()* method to create a boolean indexer.

Binning data with cut and qcut

cut() and qcut() both return a Categorical object, and the bins they create are stored as an IntervalIndex in its .categories attribute.

cut () also accepts an <code>IntervalIndex</code> for its <code>bins</code> argument, which enables a useful pandas idiom. First, We call cut () with some data and <code>bins</code> set to a fixed number, to generate the bins. Then, we pass the values of . <code>categories</code> as the <code>bins</code> argument in subsequent calls to cut (), supplying new data which will be binned into the same bins.

```
In [206]: pd.cut([0, 3, 5, 1], bins=c.categories)
Out[206]:
[(-0.003, 1.5], (1.5, 3.0], NaN, (-0.003, 1.5]]
Categories (2, interval[float64]): [(-0.003, 1.5] < (1.5, 3.0]]</pre>
```

Any value which falls outside all bins will be assigned a NaN value.

Generating ranges of intervals

If we need intervals on a regular frequency, we can use the <code>interval_range()</code> function to create an <code>IntervalIndex</code> using various combinations of start, end, and periods. The default frequency for <code>interval_range</code> is a 1 for numeric intervals, and calendar day for datetime-like intervals:

The freq parameter can used to specify non-default frequencies, and can utilize a variety of *frequency aliases* with datetime-like intervals:

```
In [210]: pd.interval_range(start=0, periods=5, freq=1.5)
Out [210]:
IntervalIndex([(0.0, 1.5], (1.5, 3.0], (3.0, 4.5], (4.5, 6.0], (6.0, 7.5]],
              closed='right',
              dtype='interval[float64]')
In [211]: pd.interval_range(start=pd.Timestamp('2017-01-01'), periods=4, freq='W')
Out [211]:
IntervalIndex([(2017-01-01, 2017-01-08], (2017-01-08, 2017-01-15], (2017-01-15, 2017-
\hookrightarrow01-22], (2017-01-22, 2017-01-29]],
              closed='right',
              dtype='interval[datetime64[ns]]')
In [212]: pd.interval_range(start=pd.Timedelta('0 days'), periods=3, freq='9H')
Out [212]:
IntervalIndex([(0 days 00:00:00, 0 days 09:00:00], (0 days 09:00:00, 0 days 18:00:00],
→ (0 days 18:00:00, 1 days 03:00:00]],
              closed='right',
              dtype='interval[timedelta64[ns]]')
```

Additionally, the closed parameter can be used to specify which side(s) the intervals are closed on. Intervals are closed on the right side by default.

New in version 0.23.0.

Specifying start, end, and periods will generate a range of evenly spaced intervals from start to end inclusively, with periods number of elements in the resulting IntervalIndex:

2.3.6 Miscellaneous indexing FAQ

Integer indexing

Label-based indexing with integer axis labels is a thorny topic. It has been discussed heavily on mailing lists and among various members of the scientific Python community. In pandas, our general viewpoint is that labels matter more than integer locations. Therefore, with an integer axis index *only* label-based indexing is possible with the standard tools like .loc. The following code will generate exceptions:

```
In [217]: s = pd.Series(range(5))
In [218]: s[-1]
KeyError
                                          Traceback (most recent call last)
<ipython-input-218-76c3dce40054> in <module>
----> 1 s[-1]
/pandas-release/pandas/pandas/core/series.py in __getitem__(self, key)
                key = com.apply_if_callable(key, self)
    870
                try:
--> 871
                    result = self.index.get_value(self, key)
    872
    873
                    if not is_scalar(result):
/pandas-release/pandas/pandas/core/indexes/base.py in get_value(self, series, key)
                k = self._convert_scalar_indexer(k, kind="getitem")
  4403
   4404
                try:
-> 4405
                    return self._engine.get_value(s, k, tz=getattr(series.dtype, "tz",
→ None))
  4406
                except KeyError as e1:
   4407
                    if len(self) > 0 and (self.holds_integer() or self.is_boolean()):
/pandas-release/pandas/pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_
→value()
/pandas-release/pandas/pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_
→value()
/pandas-release/pandas/pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_
```

```
/pandas-release/pandas/pandas/_libs/hashtable_class_helper.pxi in pandas._libs.
⇔hashtable.Int64HashTable.get_item()
/pandas-release/pandas/pandas/_libs/hashtable_class_helper.pxi in pandas._libs.
→hashtable.Int64HashTable.get_item()
KeyError: -1
In [219]: df = pd.DataFrame(np.random.randn(5, 4))
In [220]: df
Out [220]:
0 -0.130121 -0.476046 0.759104 0.213379
1 -0.082641 0.448008 0.656420 -1.051443
2 0.594956 -0.151360 -0.069303 1.221431
3 -0.182832 0.791235 0.042745 2.069775
4 1.446552 0.019814 -1.389212 -0.702312
In [221]: df.loc[-2:]
Out [221]:
                             2
                   1
0 -0.130121 -0.476046 0.759104 0.213379
1 -0.082641 0.448008 0.656420 -1.051443
2 0.594956 -0.151360 -0.069303 1.221431
3 -0.182832 0.791235 0.042745 2.069775
4 1.446552 0.019814 -1.389212 -0.702312
```

This deliberate decision was made to prevent ambiguities and subtle bugs (many users reported finding bugs when the API change was made to stop "falling back" on position-based indexing).

Non-monotonic indexes require exact matches

If the index of a Series or DataFrame is monotonically increasing or decreasing, then the bounds of a label-based slice can be outside the range of the index, much like slice indexing a normal Python list. Monotonicity of an index can be tested with the <code>is monotonic increasing()</code> and <code>is monotonic decreasing()</code> attributes.

```
Out[225]:
Empty DataFrame
Columns: [data]
Index: []
```

On the other hand, if the index is not monotonic, then both slice bounds must be *unique* members of the index.

```
# 0 is not in the index
In [9]: df.loc[0:4, :]
KeyError: 0

# 3 is not a unique label
In [11]: df.loc[2:3, :]
KeyError: 'Cannot get right slice bound for non-unique label: 3'
```

Index.is_monotonic_increasing and Index.is_monotonic_decreasing only check that an index is weakly monotonic. To check for strict monotonicity, you can combine one of those with the <code>is_unique()</code> attribute.

```
In [229]: weakly_monotonic = pd.Index(['a', 'b', 'c', 'c'])
In [230]: weakly_monotonic
Out[230]: Index(['a', 'b', 'c', 'c'], dtype='object')
In [231]: weakly_monotonic.is_monotonic_increasing
Out[231]: True
In [232]: weakly_monotonic.is_monotonic_increasing & weakly_monotonic.is_unique
Out[232]: False
```

Endpoints are inclusive

Compared with standard Python sequence slicing in which the slice endpoint is not inclusive, label-based slicing in pandas **is inclusive**. The primary reason for this is that it is often not possible to easily determine the "successor" or next element after a particular label in an index. For example, consider the following Series:

```
In [233]: s = pd.Series(np.random.randn(6), index=list('abcdef'))
In [234]: s
Out[234]:
a      0.301379
b      1.240445
c      -0.846068
d      -0.043312
e      -1.658747
f      -0.819549
dtype: float64
```

Suppose we wished to slice from c to e, using integers this would be accomplished as such:

```
In [235]: s[2:5]
Out [235]:
c -0.846068
d -0.043312
e -1.658747
dtype: float64
```

However, if you only had c and e, determining the next element in the index can be somewhat complicated. For example, the following does not work:

```
s.loc['c':'e' + 1]
```

A very common use case is to limit a time series to start and end at two specific dates. To enable this, we made the design choice to make label-based slicing include both endpoints:

```
In [236]: s.loc['c':'e']
Out[236]:
c    -0.846068
d    -0.043312
e    -1.658747
dtype: float64
```

This is most definitely a "practicality beats purity" sort of thing, but it is something to watch out for if you expect label-based slicing to behave exactly in the way that standard Python integer slicing works.

Indexing potentially changes underlying Series dtype

The different indexing operation can potentially change the dtype of a Series.

```
In [237]: series1 = pd.Series([1, 2, 3])
In [238]: series1.dtype
Out[238]: dtype('int64')
In [239]: res = series1.reindex([0, 4])
```

```
In [240]: res.dtype
Out[240]: dtype('float64')

In [241]: res
Out[241]:
0    1.0
4    NaN
dtype: float64
```

```
In [242]: series2 = pd.Series([True])
In [243]: series2.dtype
Out[243]: dtype('bool')
In [244]: res = series2.reindex_like(series1)

In [245]: res.dtype
Out[245]: dtype('0')

In [246]: res
Out[246]:
0     True
1     NaN
2     NaN
dtype: object
```

This is because the (re)indexing operations above silently inserts NaNs and the dtype changes accordingly. This can cause some issues when using numpy ufuncs such as numpy.logical_and.

See the this old issue for a more detailed discussion.

2.4 Merge, join, and concatenate

pandas provides various facilities for easily combining together Series or DataFrame with various kinds of set logic for the indexes and relational algebra functionality in the case of join / merge-type operations.

2.4.1 Concatenating objects

The <code>concat()</code> function (in the main pandas namespace) does all of the heavy lifting of performing concatenation operations along an axis while performing optional set logic (union or intersection) of the indexes (if any) on the other axes. Note that I say "if any" because there is only a single possible axis of concatenation for Series.

Before diving into all of the details of concat and what it can do, here is a simple example:

```
'C': ['C4', 'C5', 'C6', 'C7'],
   . . . :
                              'D': ['D4', 'D5', 'D6', 'D7']},
   . . . :
                             index=[4, 5, 6, 7])
   . . . :
   . . . :
In [3]: df3 = pd.DataFrame({'A': ['A8', 'A9', 'A10', 'A11'],
                              'B': ['B8', 'B9', 'B10', 'B11'],
                              'C': ['C8', 'C9', 'C10', 'C11'],
                              'D': ['D8', 'D9', 'D10', 'D11']},
   . . . :
                             index=[8, 9, 10, 11])
   . . . :
In [4]: frames = [df1, df2, df3]
In [5]: result = pd.concat(frames)
```

		df1					Result		
	Α	В	С	D		_	_	_	
0	A0	В0	α	D0		Α	В	C	D
1	A1	B1	Cl	D1	0	A0	В0	8	D0
2	A2	B2	C2	D2	1	Al	B1	C1	D1
3	A3	В3	СЗ	D3	2	A2	B2	(2	D2
		df2						_	
	Α	В	С	D	3	A3	B3	СЗ	D3
4	A4	B4	C4	D4	4	A4	B4	C4	D4
5	A5	B5	C5	D5	5	A5	B5	C5	D5
6	Аб	B6	C6	D6	6	A6	В6	C6	D6
7	A7	B7	C7	D7	7	A7	B7	C7	D7
		df3			8	A8	B8	~	DO
	Α	В	С	D	8	AB	86	C8	DB
8	A8	B8	C8	DB	9	A9	B9	C9	D9
9	A9	B9	C9	D9	10	A10	B10	C10	D10
10	A10	B10	C10	D10	11	A11	B11	C11	D11
11	A11	B11	C11	D11					

Like its sibling function on ndarrays, numpy.concatenate, pandas.concat takes a list or dict of homogeneously-typed objects and concatenates them with some configurable handling of "what to do with the other axes":

- objs: a sequence or mapping of Series or DataFrame objects. If a dict is passed, the sorted keys will be used as the *keys* argument, unless it is passed, in which case the values will be selected (see below). Any None objects will be dropped silently unless they are all None in which case a ValueError will be raised.
- axis: $\{0, 1, ...\}$, default 0. The axis to concatenate along.

- join: {'inner', 'outer'}, default 'outer'. How to handle indexes on other axis(es). Outer for union and inner for intersection.
- ignore_index: boolean, default False. If True, do not use the index values on the concatenation axis. The resulting axis will be labeled 0, ..., n 1. This is useful if you are concatenating objects where the concatenation axis does not have meaningful indexing information. Note the index values on the other axes are still respected in the join.
- keys: sequence, default None. Construct hierarchical index using the passed keys as the outermost level. If multiple levels passed, should contain tuples.
- levels: list of sequences, default None. Specific levels (unique values) to use for constructing a MultiIndex. Otherwise they will be inferred from the keys.
- names: list, default None. Names for the levels in the resulting hierarchical index.
- verify_integrity: boolean, default False. Check whether the new concatenated axis contains duplicates. This can be very expensive relative to the actual data concatenation.
- copy: boolean, default True. If False, do not copy data unnecessarily.

Without a little bit of context many of these arguments don't make much sense. Let's revisit the above example. Suppose we wanted to associate specific keys with each of the pieces of the chopped up DataFrame. We can do this using the keys argument:

		df1			Result					
	Α	В	С	D						
0	A0	B0	α	D0			Α	В	U	D
1	A1	B1	C1	D1	×	0	AD	В0	8	D0
2	A2	B2	C2	D2	×	1	A1	B1.	а	D1
3	A3	В3	СЗ	D3	×	2	A2	B2	Q	D2
		df2				_			_	
	Α	В	С	D	×	3	A3	B3	В	D3
4	A4	B4	C4	D4	У	4	A4	B4	C4	D4
5	A5	B5	C5	D5	У	5	A5	B5	0	D5
6	Аб	B6	C 6	D6	У	6	Aß	B6	C6	D6
7	A7	B7	C7	D7	У	7	A7	B7	a	D7
		df3							_	
	Α	В	С	D	z	8	AB	B8	СВ	D8
8	A8	B8	C8	DB	z	9	A9	B9	Ø	D9
9	A9	B9	C9	D9	z	10	A10	B10	Пo	D10
10	A10	B10	C10	D10	z	11	A11	B11	G1	D11
11	A11	B11	C11	D11						

As you can see (if you've read the rest of the documentation), the resulting object's index has a *hierarchical index*. This means that we can now select out each chunk by key:

```
In [7]: result.loc['y']
Out[7]:

A B C D

4 A4 B4 C4 D4

5 A5 B5 C5 D5

6 A6 B6 C6 D6

7 A7 B7 C7 D7
```

It's not a stretch to see how this can be very useful. More detail on this functionality below.

Note: It is worth noting that <code>concat()</code> (and therefore append()) makes a full copy of the data, and that constantly reusing this function can create a significant performance hit. If you need to use the operation over several datasets, use a list comprehension.

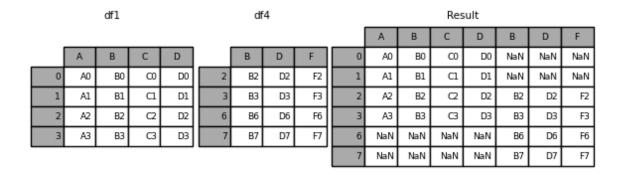
```
frames = [ process_your_file(f) for f in files ]
result = pd.concat(frames)
```

Set logic on the other axes

When gluing together multiple DataFrames, you have a choice of how to handle the other axes (other than the one being concatenated). This can be done in the following two ways:

- Take the union of them all, join='outer'. This is the default option as it results in zero information loss.
- Take the intersection, join='inner'.

Here is an example of each of these methods. First, the default join='outer' behavior:

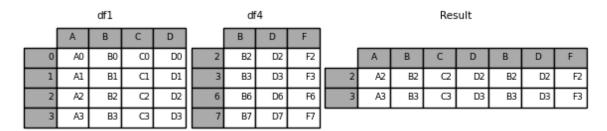


Warning: Changed in version 0.23.0.

The default behavior with join='outer' is to sort the other axis (columns in this case). In a future version of pandas, the default will be to not sort. We specified sort=False to opt in to the new behavior now.

Here is the same thing with join='inner':

```
In [10]: result = pd.concat([df1, df4], axis=1, join='inner')
```

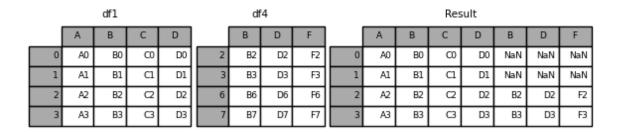


Lastly, suppose we just wanted to reuse the *exact index* from the original DataFrame:

```
In [11]: result = pd.concat([df1, df4], axis=1).reindex(df1.index)
```

Similarly, we could index before the concatenation:

```
In [12]: pd.concat([df1, df4.reindex(df1.index)], axis=1)
Out[12]:
   Α
      В
          С
             D
                  В
                       D
                           F
  AO BO CO DO NaN NaN NaN
  Α1
      В1
         C1 D1
                NaN NaN
                          NaN
  A2
      В2
         C2
             D2
                 В2
                     D2
                           F2
  A3
     B3 C3 D3
                 вЗ
```



Concatenating using append

A useful shortcut to <code>concat()</code> are the <code>append()</code> instance methods on <code>Series</code> and <code>DataFrame</code>. These methods actually predated <code>concat</code>. They concatenate along <code>axis=0</code>, namely the index:

```
In [13]: result = df1.append(df2)
```

			df1			Result				
		Α	В	С	D		Α	В	С	D
	0	A0	B0	α	D0					
	1	A1	B1	C1	D1	0	A0	B0	ω	D0
	2	A2	B2	C2	D2	1	A1	B1	C1	D1
	3	A3	В3	СЗ	D3	2	A2	B2	C2	D2
	df2						A3	В3	СЗ	D3
_		Α	В	С	D	4	A4	B4	C4	D4
	4	A4	B4	C4	D4	5	A5	B5	C5	D5
	5	A5	B5	C5	D5	-	4.5	D.C.		DE
	6	A6	B6	O6	D6	6	Аб	B6	C6	D6
	7	A7	B7	C7	D7	7	A7	В7	C7	D7

In the case of ${\tt DataFrame}$, the indexes must be disjoint but the columns do not need to be:

```
In [14]: result = df1.append(df4, sort=False)
```

			df1			Result						
_		Α	В	С	D		Α	В	С	D	F	
Г	0	A0	BO	0	0 D0							
t	1	Al	B1	c	1 D1	0	A0	BO	α	D0	NaN	
t	2	A2	B2	-	_	1	A1	B1	Cl	D1	NaN	
Ì	3	A3	В3	C	3 D3	2	A2	B2	(2	D2	NaN	
	df4						A3	В3	СЗ	D3	NaN	
_		B D		F	2	NaN	B2	NaN	D2	F2		
ı	2	1	B2	D2	F2	3	NaN	В3	NaN	D3	F3	
Γ	3	3	B3	D3	F3			D.C.		D.C.		
t	6	5	36	D6	F6	6	NaN	B6	NaN	D6	F6	
H		_	B7	D7	F7	7	NaN	В7	NaN	D7	F7	
ш	,		٠, ١	0,	l ''							

append may take multiple objects to concatenate:

```
In [15]: result = df1.append([df2, df3])
```

		df1				Result			
	Α	В	С	D		_	_		_
0	A0	В0	8	D0		Α	В	С	D
1	A1	B1	CI	D1	0	A0	B0	α	D0
2	A2	B2	C2	D2	1	Al	B1	C1	D1
3	A3	В3	СЗ	D3	2	A2	B2	C2	D2
		df2							
	Α	В	С	D	3	A3	B3	СЗ	D3
4	A4	B4	C4	D4	4	A4	B4	C4	D4
5	A5	B5	C5	D5	5	A5	B5	C5	D5
6	A6	B6	C6	D6	6	A6	B6	C6	D6
7	A7	B7	C7	D7	7	A7	B7	C7	D7
		df3			_	40			
	Α	В	С	D	8	A8	B8	C8	DB
8	A8	B8	C8	DB	9	A9	B9	8	D9
9	A9	B9	C9	D9	10	A10	B10	C10	D10
10	A10	B10	C10	D10	11	A11	B11	C11	D11
11	A11	B11	C11	D11					

Note: Unlike the append() method, which appends to the original list and returns None, append() here **does not** modify df1 and returns its copy with df2 appended.

Ignoring indexes on the concatenation axis

For DataFrame objects which don't have a meaningful index, you may wish to append them and ignore the fact that they may have overlapping indexes. To do this, use the <code>ignore_index</code> argument:

```
In [16]: result = pd.concat([df1, df4], ignore_index=True, sort=False)
```