

(continued from previous page)

```

.....:                                     '20160525 13:30:00.048',
.....:                                     '20160525 13:30:00.048',
.....:                                     '20160525 13:30:00.048'])),
.....:     'ticker': ['MSFT', 'MSFT',
.....:                'GOOG', 'GOOG', 'AAPL'],
.....:     'price': [51.95, 51.95,
.....:                720.77, 720.92, 98.00],
.....:     'quantity': [75, 155,
.....:                  100, 100, 100]},
.....:     columns=['time', 'ticker', 'price', 'quantity'])
.....:
In [134]: quotes = pd.DataFrame({
.....:     'time': pd.to_datetime(['20160525 13:30:00.023',
.....:                               '20160525 13:30:00.023',
.....:                               '20160525 13:30:00.030',
.....:                               '20160525 13:30:00.041',
.....:                               '20160525 13:30:00.048',
.....:                               '20160525 13:30:00.049',
.....:                               '20160525 13:30:00.072',
.....:                               '20160525 13:30:00.075'])),
.....:     'ticker': ['GOOG', 'MSFT', 'MSFT',
.....:                'MSFT', 'GOOG', 'AAPL', 'GOOG',
.....:                'MSFT'],
.....:     'bid': [720.50, 51.95, 51.97, 51.99,
.....:             720.50, 97.99, 720.50, 52.01],
.....:     'ask': [720.93, 51.96, 51.98, 52.00,
.....:             720.93, 98.01, 720.88, 52.03]},
.....:     columns=['time', 'ticker', 'bid', 'ask'])
.....:

```

In [135]: trades

Out[135]:

	time	ticker	price	quantity
0	2016-05-25 13:30:00.023	MSFT	51.95	75
1	2016-05-25 13:30:00.038	MSFT	51.95	155
2	2016-05-25 13:30:00.048	GOOG	720.77	100
3	2016-05-25 13:30:00.048	GOOG	720.92	100
4	2016-05-25 13:30:00.048	AAPL	98.00	100

In [136]: quotes

Out[136]:

	time	ticker	bid	ask
0	2016-05-25 13:30:00.023	GOOG	720.50	720.93
1	2016-05-25 13:30:00.023	MSFT	51.95	51.96
2	2016-05-25 13:30:00.030	MSFT	51.97	51.98
3	2016-05-25 13:30:00.041	MSFT	51.99	52.00
4	2016-05-25 13:30:00.048	GOOG	720.50	720.93
5	2016-05-25 13:30:00.049	AAPL	97.99	98.01
6	2016-05-25 13:30:00.072	GOOG	720.50	720.88
7	2016-05-25 13:30:00.075	MSFT	52.01	52.03

By default we are taking the asof of the quotes.

```

In [137]: pd.merge_asof(trades, quotes,
.....:                   on='time',

```

(continues on next page)

(continued from previous page)

```

.....:         by='ticker')
.....:
Out [137]:

```

	time	ticker	price	quantity	bid	ask
0	2016-05-25 13:30:00.023	MSFT	51.95	75	51.95	51.96
1	2016-05-25 13:30:00.038	MSFT	51.95	155	51.97	51.98
2	2016-05-25 13:30:00.048	GOOG	720.77	100	720.50	720.93
3	2016-05-25 13:30:00.048	GOOG	720.92	100	720.50	720.93
4	2016-05-25 13:30:00.048	AAPL	98.00	100	NaN	NaN

We only asof within 2ms between the quote time and the trade time.

```

In [138]: pd.merge_asof(trades, quotes,
.....:                 on='time',
.....:                 by='ticker',
.....:                 tolerance=pd.Timedelta('2ms'))
.....:
Out [138]:

```

	time	ticker	price	quantity	bid	ask
0	2016-05-25 13:30:00.023	MSFT	51.95	75	51.95	51.96
1	2016-05-25 13:30:00.038	MSFT	51.95	155	NaN	NaN
2	2016-05-25 13:30:00.048	GOOG	720.77	100	720.50	720.93
3	2016-05-25 13:30:00.048	GOOG	720.92	100	720.50	720.93
4	2016-05-25 13:30:00.048	AAPL	98.00	100	NaN	NaN

We only asof within 10ms between the quote time and the trade time and we exclude exact matches on time. Note that though we exclude the exact matches (of the quotes), prior quotes **do** propagate to that point in time.

```

In [139]: pd.merge_asof(trades, quotes,
.....:                 on='time',
.....:                 by='ticker',
.....:                 tolerance=pd.Timedelta('10ms'),
.....:                 allow_exact_matches=False)
.....:
Out [139]:

```

	time	ticker	price	quantity	bid	ask
0	2016-05-25 13:30:00.023	MSFT	51.95	75	NaN	NaN
1	2016-05-25 13:30:00.038	MSFT	51.95	155	51.97	51.98
2	2016-05-25 13:30:00.048	GOOG	720.77	100	NaN	NaN
3	2016-05-25 13:30:00.048	GOOG	720.92	100	NaN	NaN
4	2016-05-25 13:30:00.048	AAPL	98.00	100	NaN	NaN

## 2.5 Reshaping and pivot tables

### 2.5.1 Reshaping by pivoting DataFrame objects

#### Pivot

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t

df.pivot(index='foo',  
columns='bar',  
values='baz')

bar	A	B	C
foo			
one	1	2	3
two	4	5	6

Data is often stored in so-called “stacked” or “record” format:

```
In [1]: df
Out[1]:
```

	date	variable	value
0	2000-01-03	A	0.469112
1	2000-01-04	A	-0.282863
2	2000-01-05	A	-1.509059
3	2000-01-03	B	-1.135632
4	2000-01-04	B	1.212112
5	2000-01-05	B	-0.173215
6	2000-01-03	C	0.119209
7	2000-01-04	C	-1.044236
8	2000-01-05	C	-0.861849
9	2000-01-03	D	-2.104569
10	2000-01-04	D	-0.494929
11	2000-01-05	D	1.071804

For the curious here is how the above DataFrame was created:

```
import pandas._testing as tm

tm.N = 3

def unpivot(frame):
    N, K = frame.shape
    data = {'value': frame.to_numpy().ravel('F'),
            'variable': np.asarray(frame.columns).repeat(N),
            'date': np.tile(np.asarray(frame.index), K)}
```

(continues on next page)

(continued from previous page)

```

return pd.DataFrame(data, columns=['date', 'variable', 'value'])

df = unpivot(tm.makeTimeDataFrame())

```

To select out everything for variable A we could do:

```

In [2]: df[df['variable'] == 'A']
Out[2]:
   date variable  value
0 2000-01-03      A  0.469112
1 2000-01-04      A -0.282863
2 2000-01-05      A -1.509059

```

But suppose we wish to do time series operations with the variables. A better representation would be where the columns are the unique variables and an index of dates identifies individual observations. To reshape the data into this form, we use the `DataFrame.pivot()` method (also implemented as a top level function `pivot()`):

```

In [3]: df.pivot(index='date', columns='variable', values='value')
Out[3]:
variable      A      B      C      D
date
2000-01-03  0.469112 -1.135632  0.119209 -2.104569
2000-01-04 -0.282863  1.212112 -1.044236 -0.494929
2000-01-05 -1.509059 -0.173215 -0.861849  1.071804

```

If the `values` argument is omitted, and the input `DataFrame` has more than one column of values which are not used as column or index inputs to `pivot`, then the resulting “pivoted” `DataFrame` will have *hierarchical columns* whose topmost level indicates the respective value column:

```

In [4]: df['value2'] = df['value'] * 2

In [5]: pivoted = df.pivot(index='date', columns='variable')

In [6]: pivoted
Out[6]:
           value                value2
variable      A      B      C      D      A      B      C
date
2000-01-03  0.469112 -1.135632  0.119209 -2.104569  0.938225 -2.271265  0.238417 -4.
209138
2000-01-04 -0.282863  1.212112 -1.044236 -0.494929 -0.565727  2.424224 -2.088472 -0.
989859
2000-01-05 -1.509059 -0.173215 -0.861849  1.071804 -3.018117 -0.346429 -1.723698 2.
143608

```

You can then select subsets from the pivoted `DataFrame`:

```

In [7]: pivoted['value2']
Out[7]:
variable      A      B      C      D
date
2000-01-03  0.938225 -2.271265  0.238417 -4.209138

```

(continues on next page)

(continued from previous page)

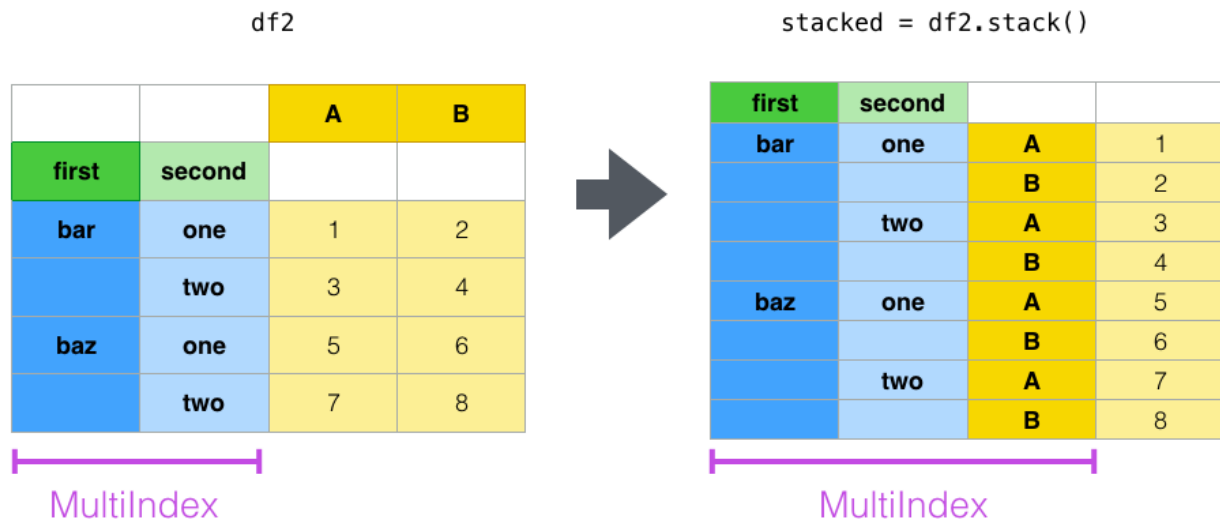
```
2000-01-04 -0.565727  2.424224 -2.088472 -0.989859
2000-01-05 -3.018117 -0.346429 -1.723698  2.143608
```

Note that this returns a view on the underlying data in the case where the data are homogeneously-typed.

**Note:** `pivot()` will error with a `ValueError: Index contains duplicate entries, cannot reshape` if the index/column pair is not unique. In this case, consider using `pivot_table()` which is a generalization of `pivot` that can handle duplicate values for one index/column pair.

## 2.5.2 Reshaping by stacking and unstacking

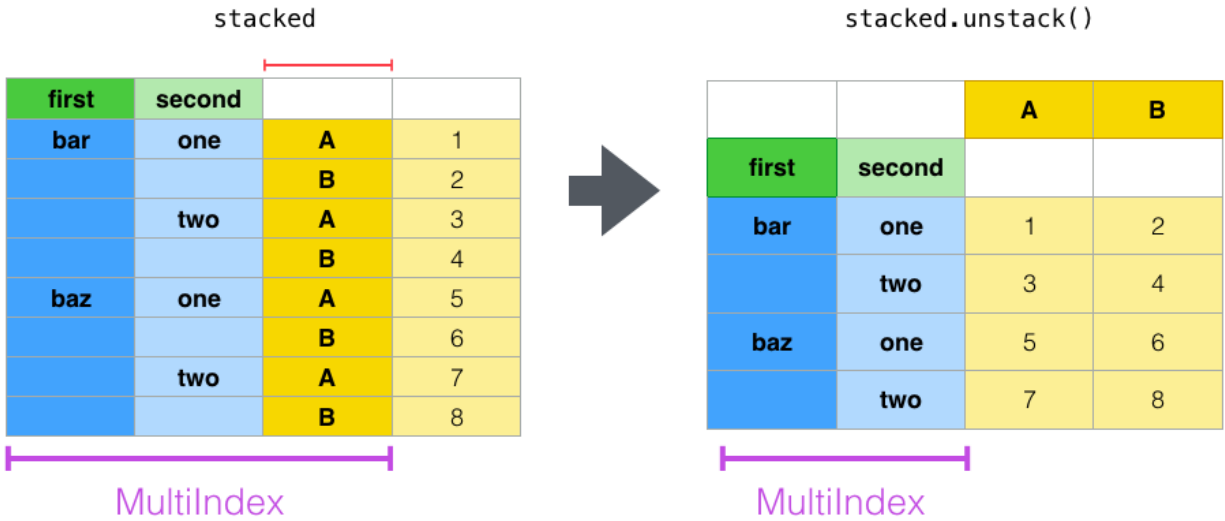
### Stack



Closely related to the `pivot()` method are the related `stack()` and `unstack()` methods available on `Series` and `DataFrame`. These methods are designed to work together with `MultiIndex` objects (see the section on *hierarchical indexing*). Here are essentially what these methods do:

- `stack`: “pivot” a level of the (possibly hierarchical) column labels, returning a `DataFrame` with an index with a new inner-most level of row labels.
- `unstack`: (inverse operation of `stack`) “pivot” a level of the (possibly hierarchical) row index to the column axis, producing a reshaped `DataFrame` with a new inner-most level of column labels.

# Unstack



The clearest way to explain is by example. Let's take a prior example data set from the hierarchical indexing section:

```
In [8]: tuples = list(zip(*[['bar', 'bar', 'baz', 'baz',
...:                        'foo', 'foo', 'qux', 'qux'],
...:                      ['one', 'two', 'one', 'two',
...:                      'one', 'two', 'one', 'two']]))
...:
In [9]: index = pd.MultiIndex.from_tuples(tuples, names=['first', 'second'])
In [10]: df = pd.DataFrame(np.random.randn(8, 2), index=index, columns=['A', 'B'])
In [11]: df2 = df[:4]
In [12]: df2
Out[12]:
```

		A	B
first	second		
bar	one	0.721555	-0.706771
	two	-1.039575	0.271860
baz	one	-0.424972	0.567020
	two	0.276232	-1.087401

The stack function “compresses” a level in the DataFrame’s columns to produce either:

- A Series, in the case of a simple column Index.
- A DataFrame, in the case of a MultiIndex in the columns.

If the columns have a MultiIndex, you can choose which level to stack. The stacked level becomes the new lowest level in a MultiIndex on the columns:

```
In [13]: stacked = df2.stack()
```

```
In [14]: stacked
```

```
Out[14]:
```

```
first second
bar   one    A    0.721555
      one    B   -0.706771
      two    A   -1.039575
      two    B    0.271860
baz   one    A   -0.424972
      one    B    0.567020
      two    A    0.276232
      two    B   -1.087401
dtype: float64
```

With a “stacked” DataFrame or Series (having a MultiIndex as the index), the inverse operation of stack is unstack, which by default unstacks the **last level**:

```
In [15]: stacked.unstack()
```

```
Out[15]:
```

```
           A           B
first second
bar   one    0.721555 -0.706771
      two   -1.039575  0.271860
baz   one   -0.424972  0.567020
      two    0.276232 -1.087401
```

```
In [16]: stacked.unstack(1)
```

```
Out[16]:
```

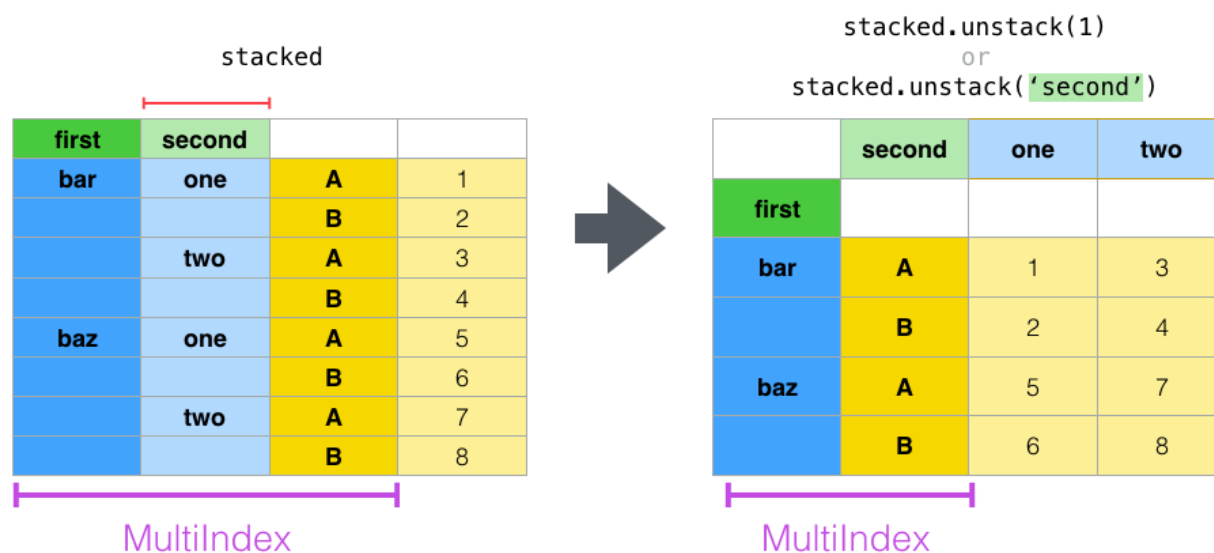
```
second      one      two
first
bar   A  0.721555 -1.039575
      B -0.706771  0.271860
baz   A -0.424972  0.276232
      B  0.567020 -1.087401
```

```
In [17]: stacked.unstack(0)
```

```
Out[17]:
```

```
first      bar      baz
second
one   A  0.721555 -0.424972
      B -0.706771  0.567020
two   A -1.039575  0.276232
      B  0.271860 -1.087401
```

## Unstack(1)

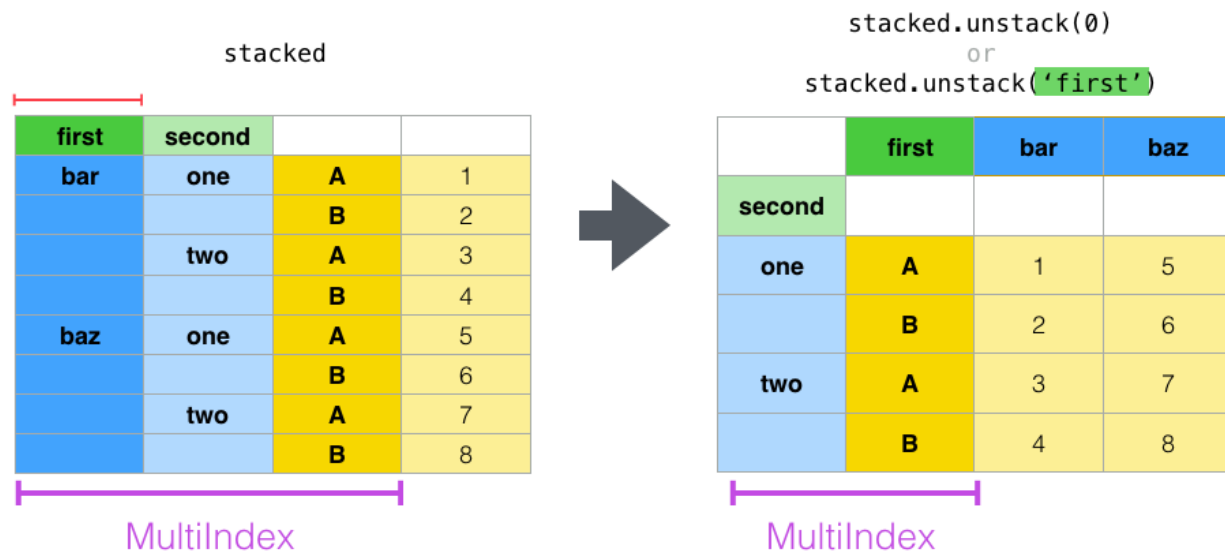


If the indexes have names, you can use the level names instead of specifying the level numbers:

```
In [18]: stacked.unstack('second')
Out[18]:
second      one      two
first
bar   A  0.721555 -1.039575
      B -0.706771  0.271860
baz   A -0.424972  0.276232
      B  0.567020 -1.087401
```



## Unstack(0)



Notice that the `stack` and `unstack` methods implicitly sort the index levels involved. Hence a call to `stack` and then `unstack`, or vice versa, will result in a **sorted** copy of the original DataFrame or Series:

```
In [19]: index = pd.MultiIndex.from_product([[2, 1], ['a', 'b']])
In [20]: df = pd.DataFrame(np.random.randn(4), index=index, columns=['A'])
In [21]: df
Out[21]:
           A
2 a -0.370647
  b -1.157892
1 a -1.344312
  b  0.844885
In [22]: all(df.unstack().stack() == df.sort_index())
Out[22]: True
```

The above code will raise a `TypeError` if the call to `sort_index` is removed.

### Multiple levels

You may also stack or unstack more than one level at a time by passing a list of levels, in which case the end result is as if each level in the list were processed individually.

```
In [23]: columns = pd.MultiIndex.from_tuples([
.....:     ('A', 'cat', 'long'), ('B', 'cat', 'long'),
.....:     ('A', 'dog', 'short'), ('B', 'dog', 'short')],
.....:     names=['exp', 'animal', 'hair_length'])
.....:
```

(continues on next page)

(continued from previous page)

```
In [24]: df = pd.DataFrame(np.random.randn(4, 4), columns=columns)
```

```
In [25]: df
```

```
Out[25]:
```

exp	A	B	A	B
animal	cat	cat	dog	dog
hair_length	long	long	short	short
0	1.075770	-0.109050	1.643563	-1.469388
1	0.357021	-0.674600	-1.776904	-0.968914
2	-1.294524	0.413738	0.276662	-0.472035
3	-0.013960	-0.362543	-0.006154	-0.923061

```
In [26]: df.stack(level=['animal', 'hair_length'])
```

```
Out[26]:
```

exp		A	B
animal	hair_length		
0	cat	long	1.075770 -0.109050
	dog	short	1.643563 -1.469388
1	cat	long	0.357021 -0.674600
	dog	short	-1.776904 -0.968914
2	cat	long	-1.294524 0.413738
	dog	short	0.276662 -0.472035
3	cat	long	-0.013960 -0.362543
	dog	short	-0.006154 -0.923061

The list of levels can contain either level names or level numbers (but not a mixture of the two).

```
# df.stack(level=['animal', 'hair_length'])
# from above is equivalent to:
```

```
In [27]: df.stack(level=[1, 2])
```

```
Out[27]:
```

exp		A	B
animal	hair_length		
0	cat	long	1.075770 -0.109050
	dog	short	1.643563 -1.469388
1	cat	long	0.357021 -0.674600
	dog	short	-1.776904 -0.968914
2	cat	long	-1.294524 0.413738
	dog	short	0.276662 -0.472035
3	cat	long	-0.013960 -0.362543
	dog	short	-0.006154 -0.923061

## Missing data

These functions are intelligent about handling missing data and do not expect each subgroup within the hierarchical index to have the same set of labels. They also can handle the index being unsorted (but you can make it sorted by calling `sort_index`, of course). Here is a more complex example:

```
In [28]: columns = pd.MultiIndex.from_tuples([('A', 'cat'), ('B', 'dog'),
.....:                                     ('B', 'cat'), ('A', 'dog')],
.....:                                     names=['exp', 'animal'])
```

```
In [29]: index = pd.MultiIndex.from_product([('bar', 'baz', 'foo', 'qux'),
.....:                                     ('one', 'two')],
```

(continues on next page)

(continued from previous page)

```

.....:                                     names=['first', 'second'])
.....:

In [30]: df = pd.DataFrame(np.random.randn(8, 4), index=index, columns=columns)

In [31]: df2 = df.iloc[[0, 1, 2, 4, 5, 7]]

In [32]: df2
Out[32]:
exp          A          B          A
animal       cat      dog      cat      dog
first second
bar   one    0.895717  0.805244 -1.206412  2.565646
      two    1.431256  1.340309 -1.170299 -0.226169
baz   one    0.410835  0.813850  0.132003 -0.827317
foo   one   -1.413681  1.607920  1.024180  0.569605
      two    0.875906 -2.211372  0.974466 -2.006747
qux   two   -1.226825  0.769804 -1.281247 -0.727707

```

As mentioned above, `stack` can be called with a `level` argument to select which level in the columns to stack:

```

In [33]: df2.stack('exp')
Out[33]:
animal          cat      dog
first second exp
bar   one    A    0.895717  2.565646
      two    B   -1.206412  0.805244
      two    A    1.431256 -0.226169
      two    B   -1.170299  1.340309
baz   one    A    0.410835 -0.827317
      two    B    0.132003  0.813850
foo   one    A   -1.413681  0.569605
      two    B    1.024180  1.607920
      two    A    0.875906 -2.006747
      two    B    0.974466 -2.211372
qux   two    A   -1.226825 -0.727707
      two    B   -1.281247  0.769804

In [34]: df2.stack('animal')
Out[34]:
exp          A          B
first second animal
bar   one    cat    0.895717 -1.206412
      two    dog    2.565646  0.805244
      two    cat    1.431256 -1.170299
      two    dog   -0.226169  1.340309
baz   one    cat    0.410835  0.132003
      two    dog   -0.827317  0.813850
foo   one    cat   -1.413681  1.024180
      two    dog    0.569605  1.607920
      two    cat    0.875906  0.974466
      two    dog   -2.006747 -2.211372
qux   two    cat   -1.226825 -1.281247
      two    dog   -0.727707  0.769804

```

Unstacking can result in missing values if subgroups do not have the same set of labels. By default, missing values will be replaced with the default fill value for that data type, `NaN` for float, `NaT` for datetimelike, etc. For integer types,

by default data will be converted to float and missing values will be set to NaN.

```
In [35]: df3 = df.iloc[[0, 1, 4, 7], [1, 2]]
```

```
In [36]: df3
```

```
Out [36]:
```

exp		B	
animal		dog	cat
	first second		
bar	one	0.805244	-1.206412
	two	1.340309	-1.170299
foo	one	1.607920	1.024180
qux	two	0.769804	-1.281247

```
In [37]: df3.unstack()
```

```
Out [37]:
```

exp		B			
animal		dog		cat	
	second	one	two	one	two
	first				
bar		0.805244	1.340309	-1.206412	-1.170299
foo		1.607920	NaN	1.024180	NaN
qux		NaN	0.769804	NaN	-1.281247

Alternatively, unstack takes an optional `fill_value` argument, for specifying the value of missing data.

```
In [38]: df3.unstack(fill_value=-1e9)
```

```
Out [38]:
```

exp		B			
animal		dog		cat	
	second	one	two	one	two
	first				
bar		8.052440e-01	1.340309e+00	-1.206412e+00	-1.170299e+00
foo		1.607920e+00	-1.000000e+09	1.024180e+00	-1.000000e+09
qux		-1.000000e+09	7.698036e-01	-1.000000e+09	-1.281247e+00

## With a MultiIndex

Unstacking when the columns are a `MultiIndex` is also careful about doing the right thing:

```
In [39]: df[:3].unstack(0)
```

```
Out [39]:
```

exp	A			B			A		
animal	cat			dog			dog		
	first	bar	baz	bar	baz	bar	baz	bar	baz
	second								
one		0.895717	0.410835	0.805244	0.81385	-1.206412	0.132003	2.565646	-0.827317
two		1.431256	NaN	1.340309	NaN	-1.170299	NaN	-0.226169	NaN

```
In [40]: df2.unstack(1)
```

```
Out [40]:
```

exp	A			B			A		
animal	cat			dog			dog		
	second	one	two	one	two	one	two	one	two
	first								
bar		0.895717	1.431256	0.805244	1.340309	-1.206412	-1.170299	2.565646	-0.226169

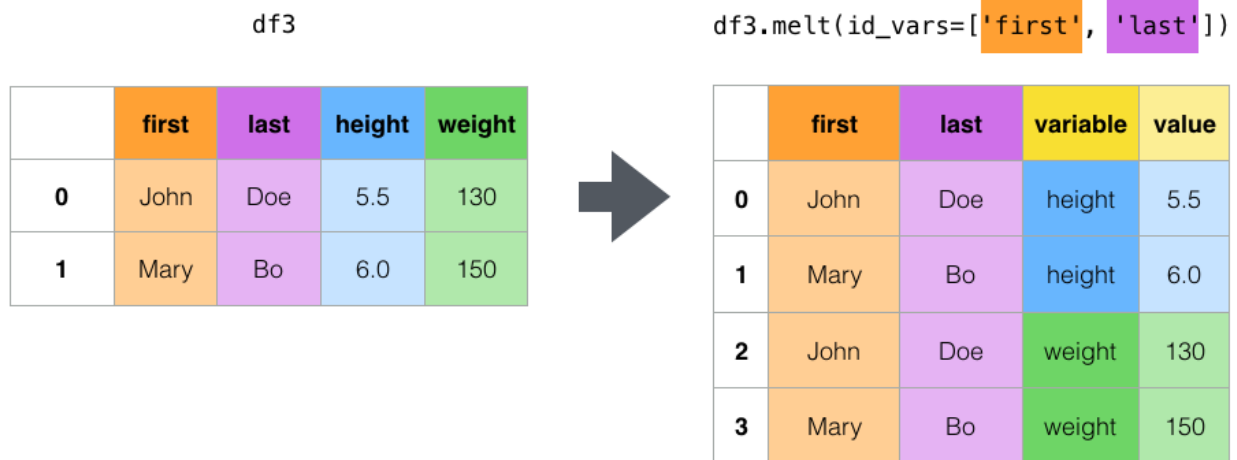
(continues on next page)

(continued from previous page)

baz	0.410835	NaN	0.813850	NaN	0.132003	NaN	-0.827317	NaN
foo	-1.413681	0.875906	1.607920	-2.211372	1.024180	0.974466	0.569605	-2.006747
qux	NaN	-1.226825	NaN	0.769804	NaN	-1.281247	NaN	-0.727707

## 2.5.3 Reshaping by Melt

### Melt



The top-level `melt()` function and the corresponding `DataFrame.melt()` are useful to massage a `DataFrame` into a format where one or more columns are *identifier variables*, while all other columns, considered *measured variables*, are “unpivoted” to the row axis, leaving just two non-identifier columns, “variable” and “value”. The names of those columns can be customized by supplying the `var_name` and `value_name` parameters.

For instance,

```
In [41]: cheese = pd.DataFrame({'first': ['John', 'Mary'],
.....:                        'last': ['Doe', 'Bo'],
.....:                        'height': [5.5, 6.0],
.....:                        'weight': [130, 150]})
.....:

In [42]: cheese
Out[42]:
   first last  height  weight
0  John  Doe     5.5     130
1  Mary  Bo      6.0     150

In [43]: cheese.melt(id_vars=['first', 'last'])
Out[43]:
   first last variable  value
0  John  Doe    height     5.5
1  Mary  Bo    height     6.0
2  John  Doe    weight    130.0
3  Mary  Bo    weight    150.0
```

(continues on next page)

(continued from previous page)

```
In [44]: cheese.melt(id_vars=['first', 'last'], var_name='quantity')
Out[44]:
```

	first	last	quantity	value
0	John	Doe	height	5.5
1	Mary	Bo	height	6.0
2	John	Doe	weight	130.0
3	Mary	Bo	weight	150.0

Another way to transform is to use the `wide_to_long()` panel data convenience function. It is less flexible than `melt()`, but more user-friendly.

```
In [45]: dft = pd.DataFrame({"A1970": {0: "a", 1: "b", 2: "c"},
.....:                      "A1980": {0: "d", 1: "e", 2: "f"},
.....:                      "B1970": {0: 2.5, 1: 1.2, 2: .7},
.....:                      "B1980": {0: 3.2, 1: 1.3, 2: .1},
.....:                      "X": dict(zip(range(3), np.random.randn(3)))
.....:                      })

In [46]: dft["id"] = dft.index

In [47]: dft
Out[47]:
```

	A1970	A1980	B1970	B1980	X	id
0	a	d	2.5	3.2	-0.121306	0
1	b	e	1.2	1.3	-0.097883	1
2	c	f	0.7	0.1	0.695775	2

```
In [48]: pd.wide_to_long(dft, ["A", "B"], i="id", j="year")
Out[48]:
```

		X	A	B
id	year			
0	1970	-0.121306	a	2.5
1	1970	-0.097883	b	1.2
2	1970	0.695775	c	0.7
0	1980	-0.121306	d	3.2
1	1980	-0.097883	e	1.3
2	1980	0.695775	f	0.1

## 2.5.4 Combining with stats and GroupBy

It should be no shock that combining pivot / stack / unstack with GroupBy and the basic Series and DataFrame statistical functions can produce some very expressive and fast data manipulations.

```
In [49]: df
Out[49]:
```

exp		A	B		A
animal		cat	dog	cat	dog
first	second				
bar	one	0.895717	0.805244	-1.206412	2.565646
	two	1.431256	1.340309	-1.170299	-0.226169
baz	one	0.410835	0.813850	0.132003	-0.827317
	two	-0.076467	-1.187678	1.130127	-1.436737
foo	one	-1.413681	1.607920	1.024180	0.569605

(continues on next page)

(continued from previous page)

```

      two      0.875906 -2.211372  0.974466 -2.006747
qux    one     -0.410001 -0.078638  0.545952 -1.219217
      two     -1.226825  0.769804 -1.281247 -0.727707

```

```
In [50]: df.stack().mean(1).unstack()
```

```
Out [50]:
```

```

animal      cat      dog
first second
bar    one     -0.155347  1.685445
      two      0.130479  0.557070
baz    one      0.271419 -0.006733
      two      0.526830 -1.312207
foo    one     -0.194750  1.088763
      two      0.925186 -2.109060
qux    one      0.067976 -0.648927
      two     -1.254036  0.021048

```

```
# same result, another way
```

```
In [51]: df.groupby(level=1, axis=1).mean()
```

```
Out [51]:
```

```

animal      cat      dog
first second
bar    one     -0.155347  1.685445
      two      0.130479  0.557070
baz    one      0.271419 -0.006733
      two      0.526830 -1.312207
foo    one     -0.194750  1.088763
      two      0.925186 -2.109060
qux    one      0.067976 -0.648927
      two     -1.254036  0.021048

```

```
In [52]: df.stack().groupby(level=1).mean()
```

```
Out [52]:
```

```

exp      A      B
second
one      0.071448  0.455513
two     -0.424186 -0.204486

```

```
In [53]: df.mean().unstack(0)
```

```
Out [53]:
```

```

exp      A      B
animal
cat      0.060843  0.018596
dog     -0.413580  0.232430

```

## 2.5.5 Pivot tables

While `pivot()` provides general purpose pivoting with various data types (strings, numerics, etc.), pandas also provides `pivot_table()` for pivoting with aggregation of numeric data.

The function `pivot_table()` can be used to create spreadsheet-style pivot tables. See the [cookbook](#) for some advanced strategies.

It takes a number of arguments:

- `data`: a `DataFrame` object.

- `values`: a column or a list of columns to aggregate.
- `index`: a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table index. If an array is passed, it is being used in the same manner as column values.
- `columns`: a column, Grouper, array which has the same length as data, or list of them. Keys to group by on the pivot table column. If an array is passed, it is being used in the same manner as column values.
- `aggfunc`: function to use for aggregation, defaulting to `numpy.mean`.

Consider a data set like this:

```
In [54]: import datetime

In [55]: df = pd.DataFrame({'A': ['one', 'one', 'two', 'three'] * 6,
.....:                    'B': ['A', 'B', 'C'] * 8,
.....:                    'C': ['foo', 'foo', 'foo', 'bar', 'bar', 'bar'] * 4,
.....:                    'D': np.random.randn(24),
.....:                    'E': np.random.randn(24),
.....:                    'F': [datetime.datetime(2013, i, 1) for i in range(1, 13)]
.....:                    + [datetime.datetime(2013, i, 15) for i in range(1, 13)]})
.....:

In [56]: df
Out[56]:
```

	A	B	C	D	E	F
0	one	A	foo	0.341734	-0.317441	2013-01-01
1	one	B	foo	0.959726	-1.236269	2013-02-01
2	two	C	foo	-1.110336	0.896171	2013-03-01
3	three	A	bar	-0.619976	-0.487602	2013-04-01
4	one	B	bar	0.149748	-0.082240	2013-05-01
..	...	..	...	...	...	...
19	three	B	foo	0.690579	-2.213588	2013-08-15
20	one	C	foo	0.995761	1.063327	2013-09-15
21	one	A	bar	2.396780	1.266143	2013-10-15
22	two	B	bar	0.014871	0.299368	2013-11-15
23	three	C	bar	3.357427	-0.863838	2013-12-15

[24 rows x 6 columns]

We can produce pivot tables from this data very easily:

```
In [57]: pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'])
Out[57]:
```

		bar	foo
one	A	1.120915	-0.514058
	B	-0.338421	0.002759
	C	-0.538846	0.699535
three	A	-1.181568	NaN
	B	NaN	0.433512
	C	0.588783	NaN
two	A	NaN	1.000985
	B	0.158248	NaN
	C	NaN	0.176180

```
In [58]: pd.pivot_table(df, values='D', index=['B'], columns=['A', 'C'], aggfunc=np.
↪sum)
Out[58]:
```

(continues on next page)



(continued from previous page)

```

A      one      three      two
C      bar      foo      bar      foo      bar      foo
B
A  2.241830 -1.028115 -2.363137      NaN      NaN  2.001971
B -0.676843  0.005518      NaN  0.867024  0.316495      NaN
C -1.077692  1.399070  1.177566      NaN      NaN  0.352360

In [59]: pd.pivot_table(df, values=['D', 'E'], index=['B'], columns=['A', 'C'],
.....:                  aggfunc=np.sum)
.....:
Out[59]:
      D      E
A      one      three      two      one
three bar      foo      bar      foo      bar      foo      bar      foo
C      bar      foo      bar      foo      bar      foo      bar      foo
B
A  2.241830 -1.028115 -2.363137      NaN      NaN  2.001971  2.786113 -0.043211  1.
922577      NaN      NaN  0.128491
B -0.676843  0.005518      NaN  0.867024  0.316495      NaN  1.368280 -1.103384
NaN -2.128743 -0.194294      NaN
C -1.077692  1.399070  1.177566      NaN      NaN  0.352360 -1.976883  1.495717 -0.
263660      NaN      NaN  0.872482

```

The result object is a DataFrame having potentially hierarchical indexes on the rows and columns. If the values column name is not given, the pivot table will include all of the data that can be aggregated in an additional level of hierarchy in the columns:

```

In [60]: pd.pivot_table(df, index=['A', 'B'], columns=['C'])
Out[60]:
      D      E
C      bar      foo      bar      foo
A      B
one  A  1.120915 -0.514058  1.393057 -0.021605
      B -0.338421  0.002759  0.684140 -0.551692
      C -0.538846  0.699535 -0.988442  0.747859
three A -1.181568      NaN  0.961289      NaN
      B      NaN  0.433512      NaN -1.064372
      C  0.588783      NaN -0.131830      NaN
two  A      NaN  1.000985      NaN  0.064245
      B  0.158248      NaN -0.097147      NaN
      C      NaN  0.176180      NaN  0.436241

```

Also, you can use `Grouper` for index and columns keywords. For detail of `Grouper`, see [Grouping with a Grouper specification](#).

```

In [61]: pd.pivot_table(df, values='D', index=pd.Grouper(freq='M', key='F'),
.....:                  columns='C')
.....:
Out[61]:
C      bar      foo
F
2013-01-31      NaN -0.514058
2013-02-28      NaN  0.002759

```

(continues on next page)

(continued from previous page)

2013-03-31	NaN	0.176180
2013-04-30	-1.181568	NaN
2013-05-31	-0.338421	NaN
2013-06-30	-0.538846	NaN
2013-07-31	NaN	1.000985
2013-08-31	NaN	0.433512
2013-09-30	NaN	0.699535
2013-10-31	1.120915	NaN
2013-11-30	0.158248	NaN
2013-12-31	0.588783	NaN

You can render a nice output of the table omitting the missing values by calling `to_string` if you wish:

```
In [62]: table = pd.pivot_table(df, index=['A', 'B'], columns=['C'])
```

```
In [63]: print(table.to_string(na_rep=''))
```

		D		E	
C		bar	foo	bar	foo
one	A	1.120915	-0.514058	1.393057	-0.021605
	B	-0.338421	0.002759	0.684140	-0.551692
	C	-0.538846	0.699535	-0.988442	0.747859
three	A	-1.181568		0.961289	
	B		0.433512		-1.064372
	C	0.588783		-0.131830	
two	A		1.000985		0.064245
	B	0.158248		-0.097147	
	C		0.176180		0.436241

Note that `pivot_table` is also available as an instance method on `DataFrame`, i.e. `DataFrame.pivot_table()`.

*DataFrame.*

## Adding margins

If you pass `margins=True` to `pivot_table`, special All columns and rows will be added with partial group aggregates across the categories on the rows and columns:

```
In [64]: df.pivot_table(index=['A', 'B'], columns='C', margins=True, aggfunc=np.std)
```

```
Out [64]:
```

		D		E	
C		bar	foo	bar	foo
one	A	1.804346	1.210272	1.569879	0.179483
	B	0.690376	1.353355	0.898998	1.083825
	C	0.273641	0.418926	0.771139	1.689271
three	A	0.794212	NaN	0.794212	2.049040
	B	NaN	0.363548	0.363548	NaN
	C	3.915454	NaN	3.915454	1.035215
two	A	NaN	0.442998	0.442998	NaN
	B	0.202765	NaN	0.202765	0.560757
	C	NaN	1.819408	1.819408	NaN
All		1.556686	0.952552	1.246608	1.250924

## 2.5.6 Cross tabulations

Use `crosstab()` to compute a cross-tabulation of two (or more) factors. By default `crosstab` computes a frequency table of the factors unless an array of values and an aggregation function are passed.

It takes a number of arguments

- `index`: array-like, values to group by in the rows.
- `columns`: array-like, values to group by in the columns.
- `values`: array-like, optional, array of values to aggregate according to the factors.
- `aggfunc`: function, optional, If no values array is passed, computes a frequency table.
- `rownames`: sequence, default `None`, must match number of row arrays passed.
- `colnames`: sequence, default `None`, if passed, must match number of column arrays passed.
- `margins`: boolean, default `False`, Add row/column margins (subtotals)
- `normalize`: boolean, {'all', 'index', 'columns'}, or {0,1}, default `False`. Normalize by dividing all values by the sum of values.

Any Series passed will have their name attributes used unless row or column names for the cross-tabulation are specified

For example:

```
In [65]: foo, bar, dull, shiny, one, two = 'foo', 'bar', 'dull', 'shiny', 'one', 'two'
In [66]: a = np.array([foo, foo, bar, bar, foo, foo], dtype=object)
In [67]: b = np.array([one, one, two, one, two, one], dtype=object)
In [68]: c = np.array([dull, dull, shiny, dull, dull, shiny], dtype=object)
In [69]: pd.crosstab(a, [b, c], rownames=['a'], colnames=['b', 'c'])
Out[69]:
b      one      two
c  dull shiny dull shiny
a
bar    1      0      0      1
foo    2      1      1      0
```

If `crosstab` receives only two Series, it will provide a frequency table.

```
In [70]: df = pd.DataFrame({'A': [1, 2, 2, 2, 2], 'B': [3, 3, 4, 4, 4],
.....:                    'C': [1, 1, np.nan, 1, 1]})
.....:
In [71]: df
Out[71]:
   A  B    C
0  1  3  1.0
1  2  3  1.0
2  2  4  NaN
3  2  4  1.0
4  2  4  1.0
In [72]: pd.crosstab(df['A'], df['B'])
```

(continues on next page)

(continued from previous page)

**Out [72]:**

```
B    3    4
A
1    1    0
2    1    3
```

Any input passed containing `Categorical` data will have **all** of its categories included in the cross-tabulation, even if the actual data does not contain any instances of a particular category.

```
In [73]: foo = pd.Categorical(['a', 'b'], categories=['a', 'b', 'c'])
```

```
In [74]: bar = pd.Categorical(['d', 'e'], categories=['d', 'e', 'f'])
```

```
In [75]: pd.crosstab(foo, bar)
```

**Out [75]:**

```
col_0  d  e
row_0
a       1  0
b       0  1
```

## Normalization

Frequency tables can also be normalized to show percentages rather than counts using the `normalize` argument:

```
In [76]: pd.crosstab(df['A'], df['B'], normalize=True)
```

**Out [76]:**

```
B      3      4
A
1  0.2  0.0
2  0.2  0.6
```

`normalize` can also normalize values within each row or within each column:

```
In [77]: pd.crosstab(df['A'], df['B'], normalize='columns')
```

**Out [77]:**

```
B      3      4
A
1  0.5  0.0
2  0.5  1.0
```

`crosstab` can also be passed a third `Series` and an aggregation function (`aggfunc`) that will be applied to the values of the third `Series` within each group defined by the first two `Series`:

```
In [78]: pd.crosstab(df['A'], df['B'], values=df['C'], aggfunc=np.sum)
```

**Out [78]:**

```
B      3      4
A
1  1.0  NaN
2  1.0  2.0
```

## Adding margins

Finally, one can also add margins or normalize this output.

```
In [79]: pd.crosstab(df['A'], df['B'], values=df['C'], aggfunc=np.sum, normalize=True,
.....:               margins=True)
.....:
Out[79]:
B      3      4    All
A
1      0.25  0.0  0.25
2      0.25  0.5  0.75
All    0.50  0.5  1.00
```

## 2.5.7 Tiling

The `cut()` function computes groupings for the values of the input array and is often used to transform continuous variables to discrete or categorical variables:

```
In [80]: ages = np.array([10, 15, 13, 12, 23, 25, 28, 59, 60])

In [81]: pd.cut(ages, bins=3)
Out[81]:
[(9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.95, 26.667], (9.
→95, 26.667], (26.667, 43.333], (43.333, 60.0], (43.333, 60.0]]
Categories (3, interval[float64]): [(9.95, 26.667] < (26.667, 43.333] < (43.333, 60.
→0]]
```

If the `bins` keyword is an integer, then equal-width bins are formed. Alternatively we can specify custom bin-edges:

```
In [82]: c = pd.cut(ages, bins=[0, 18, 35, 70])

In [83]: c
Out[83]:
[(0, 18], (0, 18], (0, 18], (0, 18], (18, 35], (18, 35], (18, 35], (35, 70], (35, 70]]
Categories (3, interval[int64]): [(0, 18] < (18, 35] < (35, 70]]
```

If the `bins` keyword is an `IntervalIndex`, then these will be used to bin the passed data.:

```
pd.cut([25, 20, 50], bins=c.categories)
```

## 2.5.8 Computing indicator / dummy variables

To convert a categorical variable into a “dummy” or “indicator” `DataFrame`, for example a column in a `DataFrame` (a `Series`) which has  $k$  distinct values, can derive a `DataFrame` containing  $k$  columns of 1s and 0s using `get_dummies()`:

```
In [84]: df = pd.DataFrame({'key': list('bbacab'), 'data1': range(6)})

In [85]: pd.get_dummies(df['key'])
Out[85]:
   a  b  c
0  0  1  0
1  0  1  0
```

(continues on next page)

(continued from previous page)

```

2  1  0  0
3  0  0  1
4  1  0  0
5  0  1  0

```

Sometimes it's useful to prefix the column names, for example when merging the result with the original DataFrame:

```
In [86]: dummies = pd.get_dummies(df['key'], prefix='key')
```

```
In [87]: dummies
```

```
Out[87]:
```

```

   key_a  key_b  key_c
0      0     1     0
1      0     1     0
2      1     0     0
3      0     0     1
4      1     0     0
5      0     1     0

```

```
In [88]: df[['data1']].join(dummies)
```

```
Out[88]:
```

```

   data1  key_a  key_b  key_c
0      0     0     1     0
1      1     0     1     0
2      2     1     0     0
3      3     0     0     1
4      4     1     0     0
5      5     0     1     0

```

This function is often used along with discretization functions like `cut`:

```
In [89]: values = np.random.randn(10)
```

```
In [90]: values
```

```
Out[90]:
```

```

array([ 0.4082, -1.0481, -0.0257, -0.9884,  0.0941,  1.2627,  1.29   ,
        0.0824, -0.0558,  0.5366])

```

```
In [91]: bins = [0, 0.2, 0.4, 0.6, 0.8, 1]
```

```
In [92]: pd.get_dummies(pd.cut(values, bins))
```

```
Out[92]:
```

```

   (0.0, 0.2]  (0.2, 0.4]  (0.4, 0.6]  (0.6, 0.8]  (0.8, 1.0]
0            0            0            1            0            0
1            0            0            0            0            0
2            0            0            0            0            0
3            0            0            0            0            0
4            1            0            0            0            0
5            0            0            0            0            0
6            0            0            0            0            0
7            1            0            0            0            0
8            0            0            0            0            0
9            0            0            1            0            0

```

See also `Series.str.get_dummies`.

`get_dummies()` also accepts a DataFrame. By default all categorical variables (categorical in the statistical sense, those with *object* or *categorical* dtype) are encoded as dummy variables.

```
In [93]: df = pd.DataFrame({'A': ['a', 'b', 'a'], 'B': ['c', 'c', 'b'],
.....:                    'C': [1, 2, 3]})
.....:
```

```
In [94]: pd.get_dummies(df)
```

```
Out[94]:
```

	C	A_a	A_b	B_b	B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

All non-object columns are included untouched in the output. You can control the columns that are encoded with the `columns` keyword.

```
In [95]: pd.get_dummies(df, columns=['A'])
```

```
Out[95]:
```

	B	C	A_a	A_b
0	c	1	1	0
1	c	2	0	1
2	b	3	1	0

Notice that the B column is still included in the output, it just hasn't been encoded. You can drop B before calling `get_dummies` if you don't want to include it in the output.

As with the `Series` version, you can pass values for the `prefix` and `prefix_sep`. By default the column name is used as the prefix, and `'_'` as the prefix separator. You can specify `prefix` and `prefix_sep` in 3 ways:

- string: Use the same value for `prefix` or `prefix_sep` for each column to be encoded.
- list: Must be the same length as the number of columns being encoded.
- dict: Mapping column name to prefix.

```
In [96]: simple = pd.get_dummies(df, prefix='new_prefix')
```

```
In [97]: simple
```

```
Out[97]:
```

	C	new_prefix_a	new_prefix_b	new_prefix_b	new_prefix_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [98]: from_list = pd.get_dummies(df, prefix=['from_A', 'from_B'])
```

```
In [99]: from_list
```

```
Out[99]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

```
In [100]: from_dict = pd.get_dummies(df, prefix={'B': 'from_B', 'A': 'from_A'})
```

```
In [101]: from_dict
```

```
Out[101]:
```

	C	from_A_a	from_A_b	from_B_b	from_B_c
0	1	1	0	0	1
1	2	0	1	0	1
2	3	1	0	1	0

Sometimes it will be useful to only keep k-1 levels of a categorical variable to avoid collinearity when feeding the result to statistical models. You can switch to this mode by turn on `drop_first`.

```
In [102]: s = pd.Series(list('abcaa'))

In [103]: pd.get_dummies(s)
Out[103]:
```

	a	b	c
0	1	0	0
1	0	1	0
2	0	0	1
3	1	0	0
4	1	0	0

```

In [104]: pd.get_dummies(s, drop_first=True)
Out[104]:
```

	b	c
0	0	0
1	1	0
2	0	1
3	0	0
4	0	0

When a column contains only one level, it will be omitted in the result.

```
In [105]: df = pd.DataFrame({'A': list('aaaaa'), 'B': list('ababc')})

In [106]: pd.get_dummies(df)
Out[106]:
```

	A_a	B_a	B_b	B_c
0	1	1	0	0
1	1	0	1	0
2	1	1	0	0
3	1	0	1	0
4	1	0	0	1

```

In [107]: pd.get_dummies(df, drop_first=True)
Out[107]:
```

	B_b	B_c
0	0	0
1	1	0
2	0	0
3	1	0
4	0	1

By default new columns will have `np.uint8` dtype. To choose another dtype, use the `dtype` argument:

```
In [108]: df = pd.DataFrame({'A': list('abc'), 'B': [1.1, 2.2, 3.3]})

In [109]: pd.get_dummies(df, dtype=bool).dtypes
Out[109]:
```

B		float64
A_a		bool
A_b		bool
A_c		bool
dtype:		object

New in version 0.23.0.



## 2.5.9 Factorizing values

To encode 1-d values as an enumerated type use `factorize()`:

```
In [110]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])

In [111]: x
Out[111]:
0      A
1      A
2     NaN
3      B
4    3.14
5     inf
dtype: object

In [112]: labels, uniques = pd.factorize(x)

In [113]: labels
Out[113]: array([ 0,  0, -1,  1,  2,  3])

In [114]: uniques
Out[114]: Index(['A', 'B', 3.14, inf], dtype='object')
```

Note that `factorize` is similar to `numpy.unique`, but differs in its handling of NaN:

---

**Note:** The following `numpy.unique` will fail under Python 3 with a `TypeError` because of an ordering bug. See also [here](#).

---

```
In [1]: x = pd.Series(['A', 'A', np.nan, 'B', 3.14, np.inf])
In [2]: pd.factorize(x, sort=True)
Out[2]:
(array([ 2,  2, -1,  3,  0,  1]),
 Index([3.14, inf, 'A', 'B'], dtype='object'))

In [3]: np.unique(x, return_inverse=True)[::-1]
Out[3]: (array([3, 3, 0, 4, 1, 2]), array([nan, 3.14, inf, 'A', 'B'], dtype=object))
```

---

**Note:** If you just want to handle one column as a categorical variable (like R's factor), you can use `df["cat_col"] = pd.Categorical(df["col"])` or `df["cat_col"] = df["col"].astype("category")`. For full docs on *Categorical*, see the *Categorical introduction* and the *API documentation*.

---

## 2.5.10 Examples

In this section, we will review frequently asked questions and examples. The column names and relevant column values are named to correspond with how this DataFrame will be pivoted in the answers below.

```
In [115]: np.random.seed([3, 1415])

In [116]: n = 20

In [117]: cols = np.array(['key', 'row', 'item', 'col'])
```

(continues on next page)